



Large Language Model Powered Conversational Agents

Yang Deng

May 13, 2024

Large Language Model Powered Conversational Systems



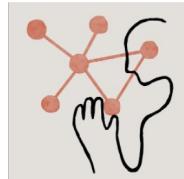
ChatGPT



Gemini



New Bing



Claude

...



Alpaca



Vicuna



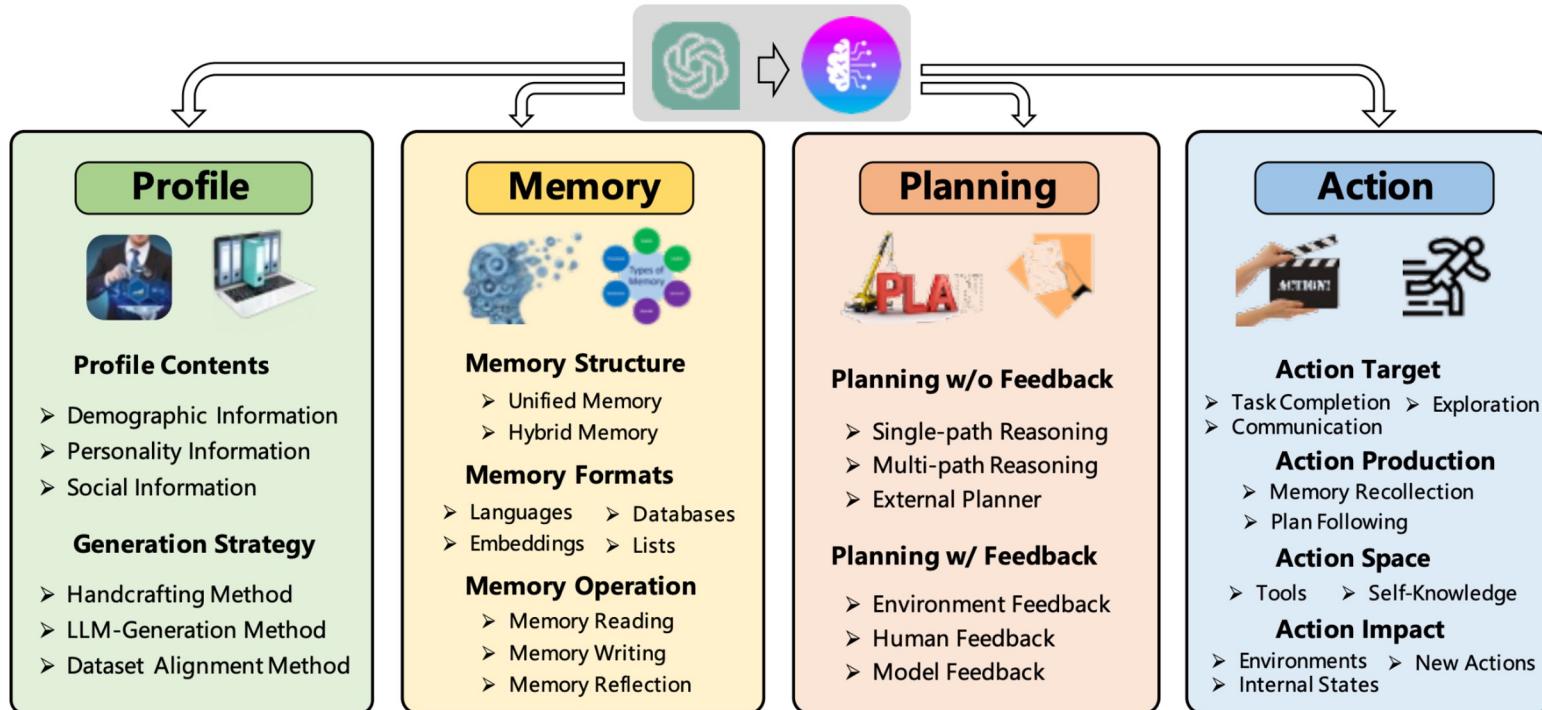
Dolly



LLaMA-Chat

Powerful capabilities of
Context Understanding
& Response Generation

LLM-powered Conversational Agents?



Overview of LLM-powered Conversational Agents



Profile

LLM-powered Conversational Agents for **User Simulation**



Memory

LLM-powered Conversational Agents for **Long-context Dialogues**



Planning

LLM-powered Conversational Agents for **Proactive Dialogues**



Action

LLM-powered Conversational Agents for **Real-world Problem Solving**

User Simulators in the Pre-LLM Era

□ User Satisfaction Estimation

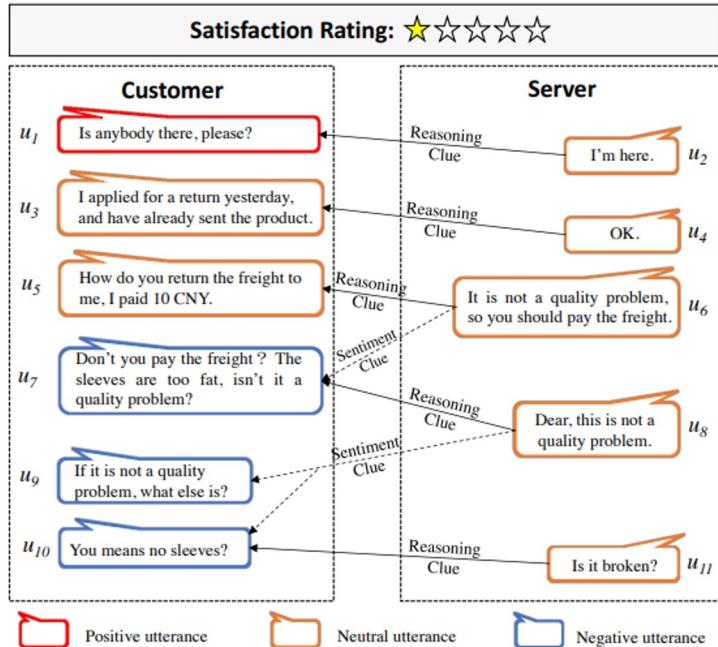
- 1) Semantic-based Estimation
- 2) Preference-based Estimation
- 3) Action-based Estimation

□ User Response Simulation

- 1) Retrieval-based User Simulators
- 2) Schema-based User Simulators
- 3) Conditioned Generation Models as User Simulators

Semantic-based User Satisfaction Estimation

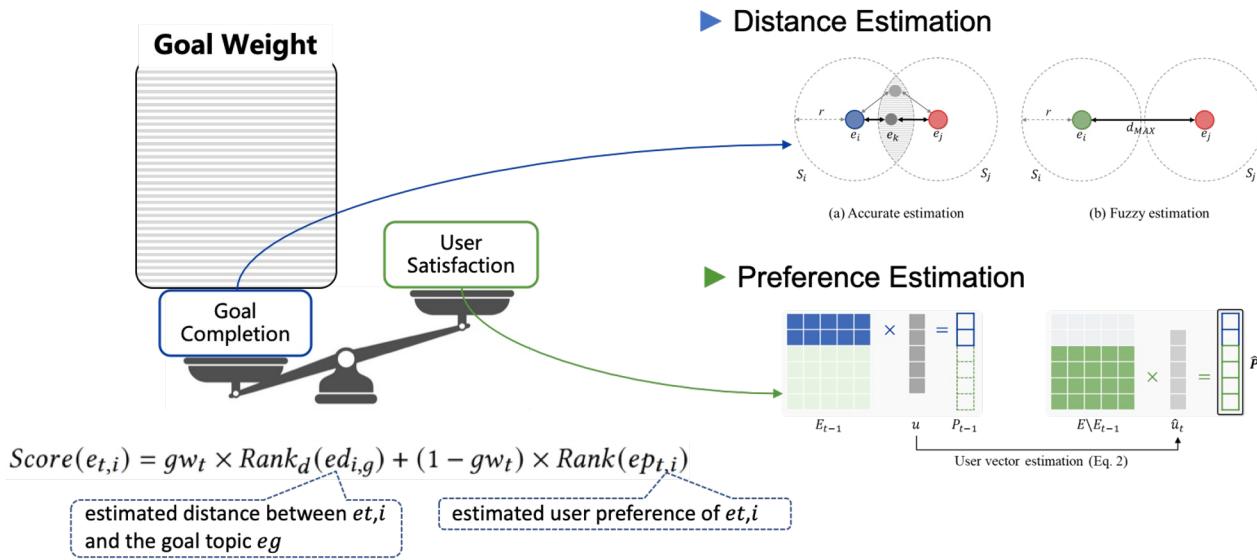
Sentiment Classification



Response Quality Assessment

Utterance	Attention weight and turn-level ratings
User: Open {restaurant booking app.}	System: Hey this is {application} $\{W_{attn} : 0.126, \hat{RQ} : 5.0, RQ : 5.0\}$
User: Can you can you uh book a table at koreana for two ?	System: Sorry, this is not curr. supported. $\{W_{attn} : 0.147, \hat{RQ} : 1.8, RQ : 2.0\}$
User: Can you book a table at the cheesecake factory for three ?	System: What time would you like to make the reservation for? $\{W_{attn} : 0.137, \hat{RQ} : 3.2, RQ : 4.0\}$
User: Three oclock	System: What date would you like the reservation for? $\{W_{attn} : 0.135, \hat{RQ} : 3.6, RQ : 4.0\}$
User: For tomorrow	System: Sorry. No restaurants could be found that match your criteria. $\{W_{attn} : 0.152, \hat{RQ} : 1.1, RQ : 2.0\}$
User: Ok	System: {No response} $\{W_{attn} : 0.153, \hat{RQ} : 1.0, RQ : 4.0\}$
User: Stop	System: {No response} $\{W_{attn} : 0.149, \hat{RQ} : 1.5, RQ : 4.0\}$

Preference-based User Satisfaction Estimation



Satisfaction is formalized as the cumulative average of users' preferences for the topics covered by the conversation:

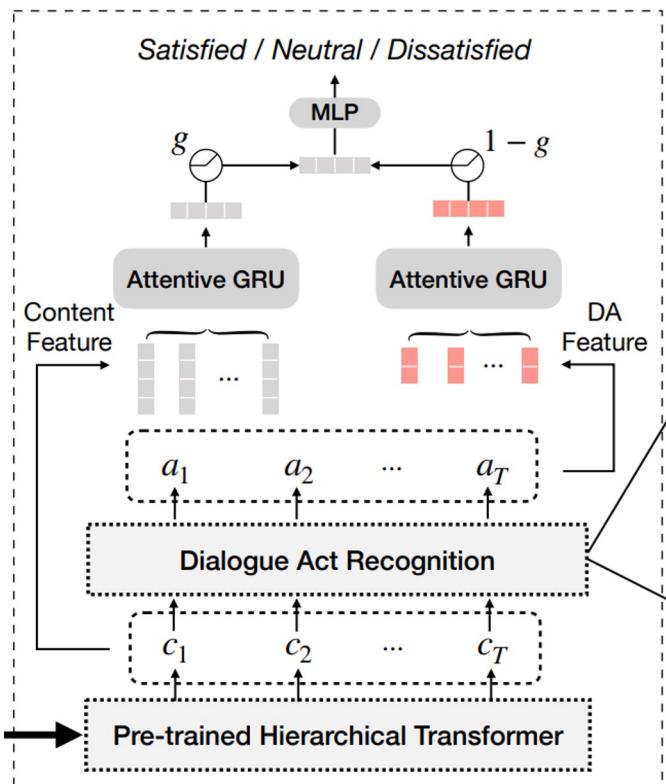
$$US_t \triangleq \frac{1}{t} \sum_{i=1}^t \frac{1}{|u_i+1|} \left(\sum_{j=1}^{|u_i|} p_{e_{i,j}} + p_{e_i^a} \right)$$

Action-based User Satisfaction Estimation

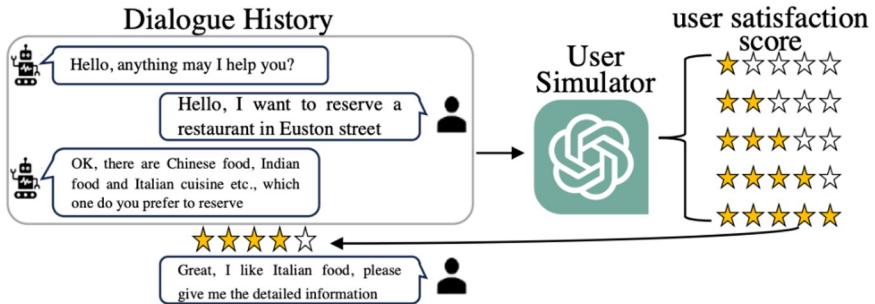
Satisfaction R		
	Is anybody there?	SAT
	Yes, what?	SGD
	The phone is hot when looking for...	DSAT
You can apply (takes long contact a rep)		SAT
	Besides me should I	MWOZ
Mobile photo electronic invc		DSAT
	Is it okay shot?	SAT
	Yes,	JDDC
	OK, I will t	DSAT

Sequence of actions:

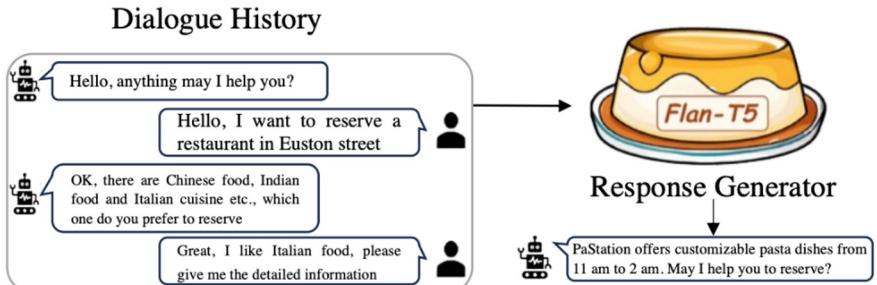
- SGD: 1. INFORM_INTENT → SELECT → AFFIRM_INTENT → AFFIRM
2. THANK_YOU → AFFIRM → THANK_YOU
3. INFORM → SELECT → INFORM_INTENT → SELECT
4. SELECT → THANK_YOU
5. AFFIRM → THANK_YOU → AFFIRM → THANK_YOU
- DSAT: 1. REQUEST → SELECT → REQUEST_ALTS → REQUEST_ALTS
2. NEGATE
3. AFFIRM → INFORM → AFFIRM → NEGATE
4. AFFIRM → AFFIRM → NEGATE
5. AFFIRM → INFORM_INTENT → INFORM → REQUEST_ALTS
- MWOZ: 1. general-thank → Restaurant-Inform → Restaurant-Request
2. Attraction-Request → Attraction-Request → general-bye
3. Attraction-Inform → Taxi-Inform → general-thank
4. general-thank → general-thank
5. general-thank → general-bye
- DSAT: 1. general-greet → Restaurant-Inform → Other → Other
2. Taxi-Inform → Taxi-Inform → Train-Inform
3. Hotel-Inform → Attraction-Request → Hotel-Inform
4. Taxi-Inform → Taxi-Inform → Taxi-Inform
5. Attraction-Request → Attraction-Request → Other → Other
- SAT: 1. Gifts for Writing Reviews → Review Viewing
2. Invoice Return&Modification → OTHER → Invoice Make-up
3. Usage Instruction → Application Instruction → OTHER
4. Processing Time of Order Cancellation → Order Resume
5. Invoice Checking → OTHER → Delivery Period
- DSAT: 1. No Record → Mail Refuse → Mail Tracking
2. Warranty&Return Policy → Unable to Apply for Insurance
3. Warranty&Return Policy → VIP → Warranty&Return Policy
4. Promotion Form → Upcoming Events → Promotion Form
5. Contact Manual Service → OTHER → Contact Manual Service



LLMs for User Satisfaction Estimation

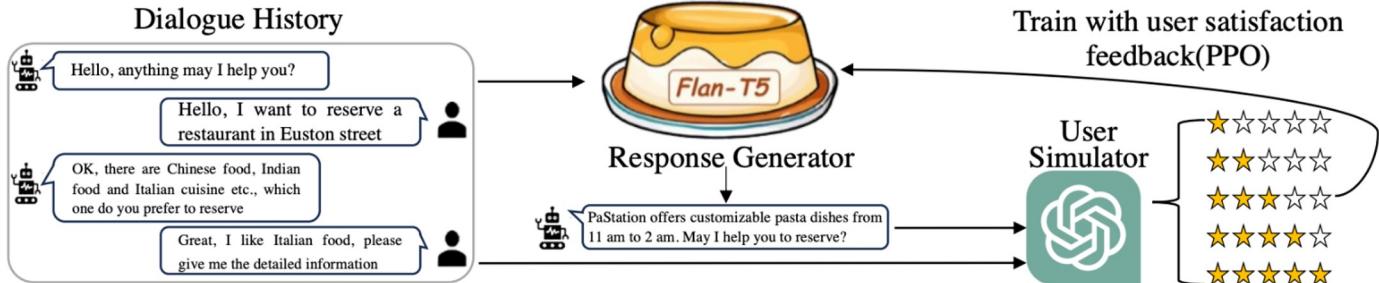


a) LLM Serve as User Simulator



b) Supervised Training of TOD Model

c) User-Guided Response Optimization



User Simulators in the Pre-LLM Era

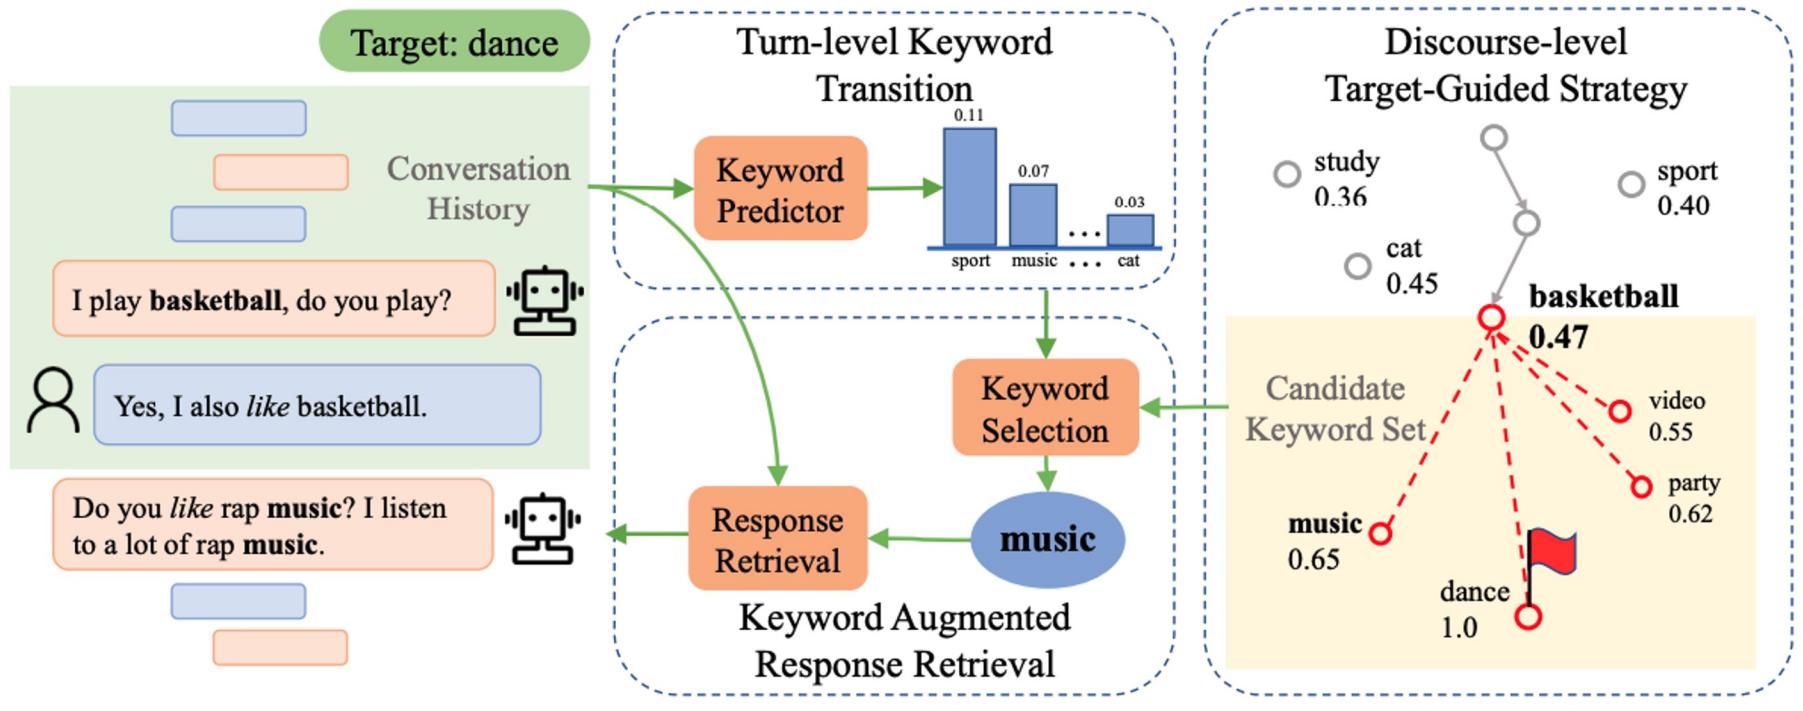
□ User Satisfaction Estimation

- 1) Semantic-based Estimation
- 2) Preference-based Estimation
- 3) Action-based Estimation

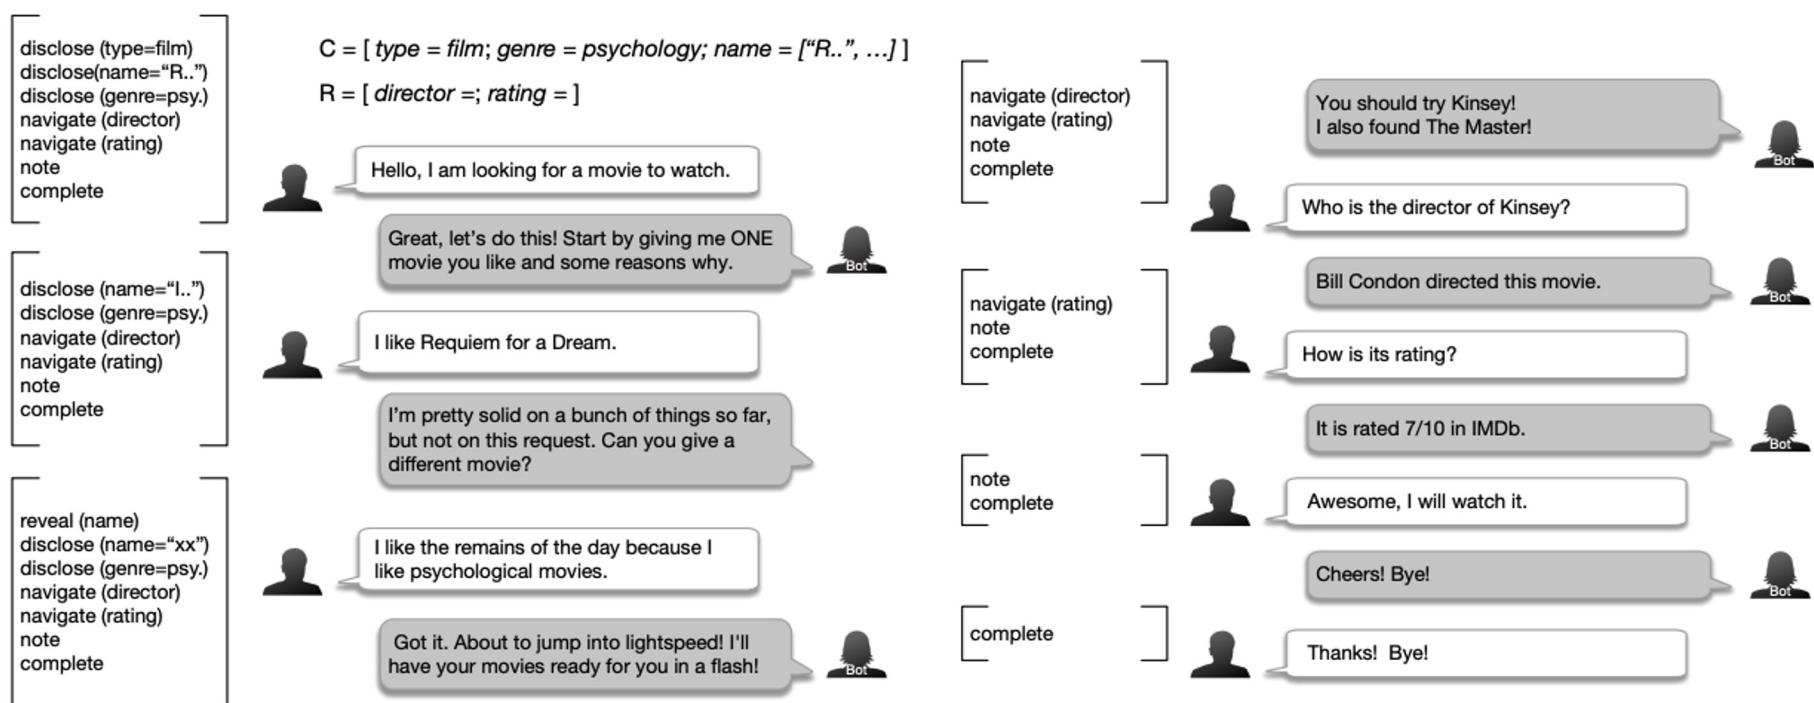
□ User Response Simulation

- 1) Retrieval-based User Simulators
- 2) Schema-based User Simulators
- 3) Conditioned Generation Models as User Simulators

Retrieval-based User Simulators

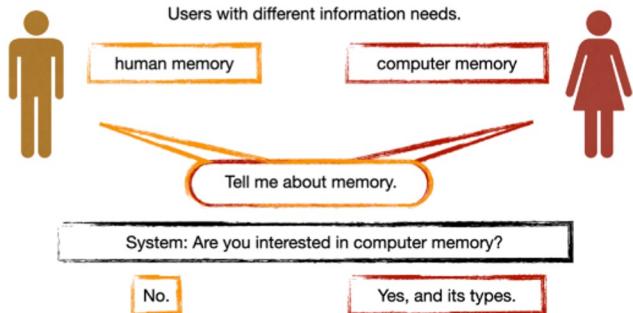
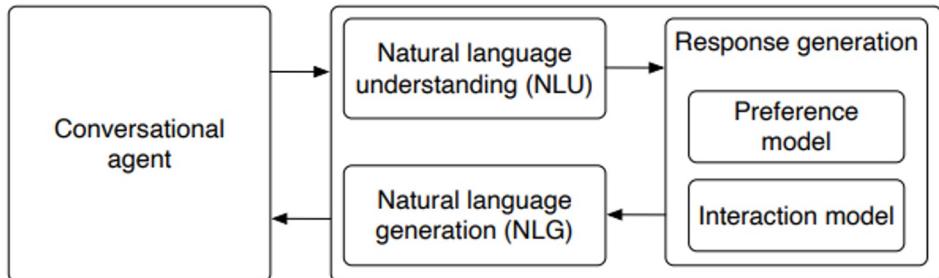


Schema-based User Simulators



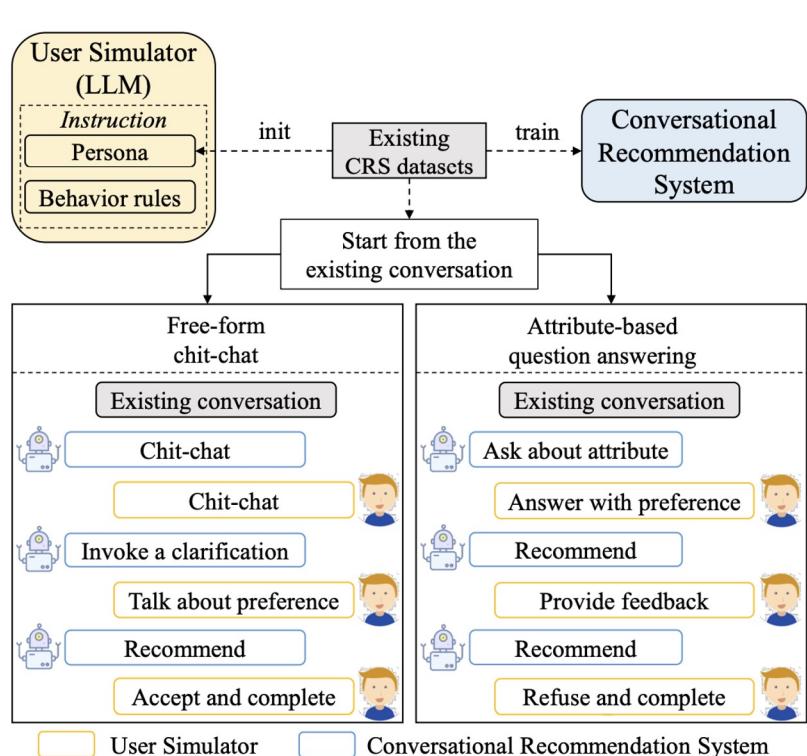
Conditional Generation Models as User Simulators

Conditioned on **user preferences** for evaluating conversational recommender systems.



Conditioned on **information needs** for evaluating conversational search systems.

LLM-powered Conversational Agents as User Simulators

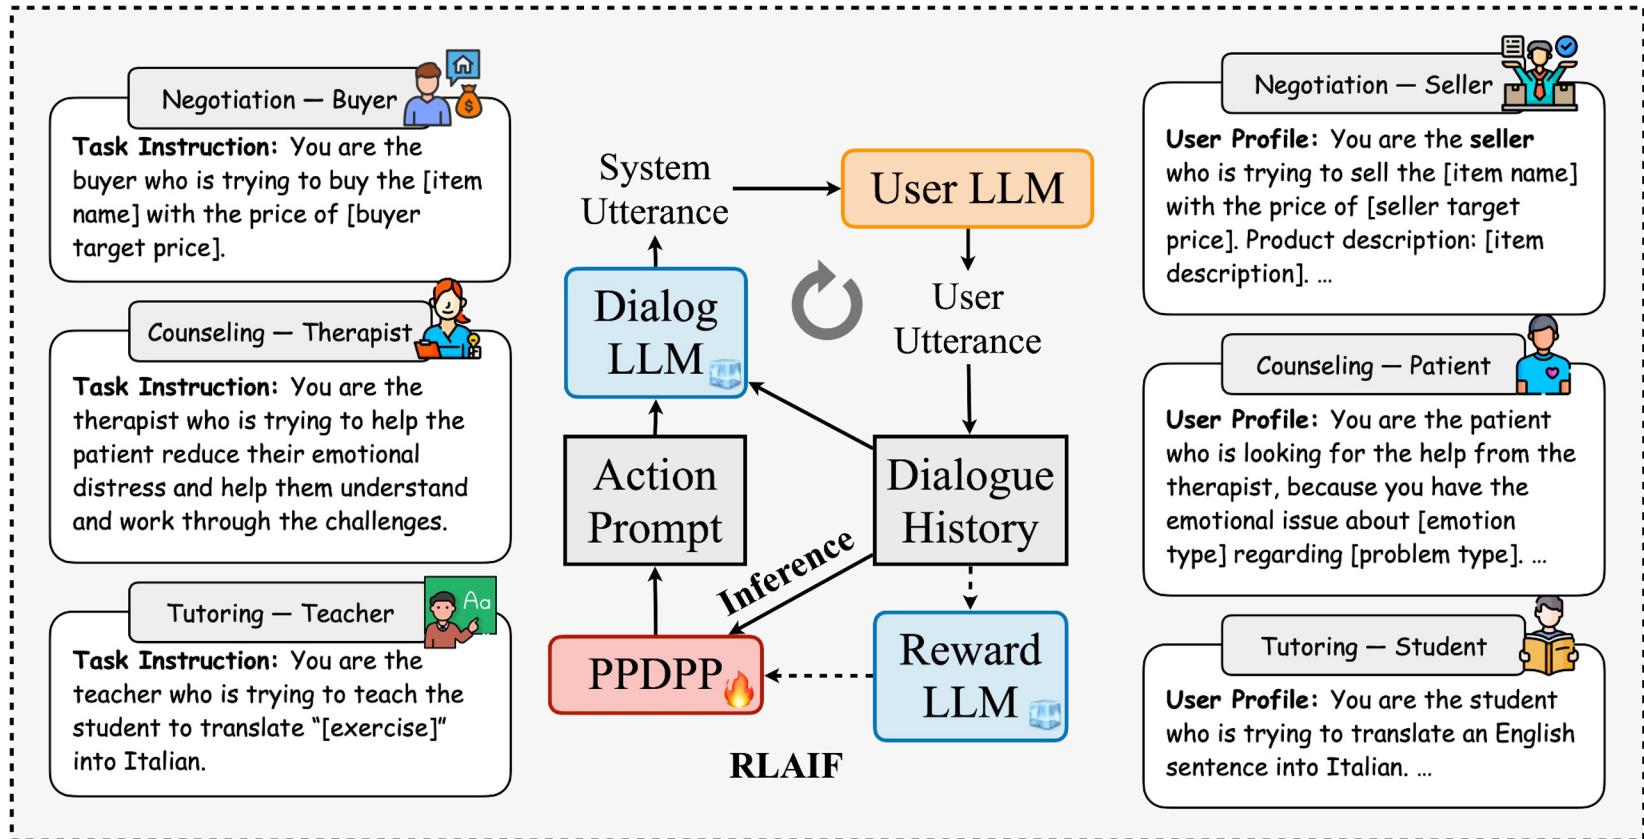


LLMs possess excellent *role-playing* capacities.

Example: Conversational Recommendation

- User Profiling / Persona:
 - *Target Items*
 - *Preferred Attributes*
- Action / Behavior Rule:
 - *Talking about preference*
 - *Providing feedback*
 - *Completing the conversation*

Role-playing Agents for Diverse Applications



Role-playing Agents for Simulating Diverse Users

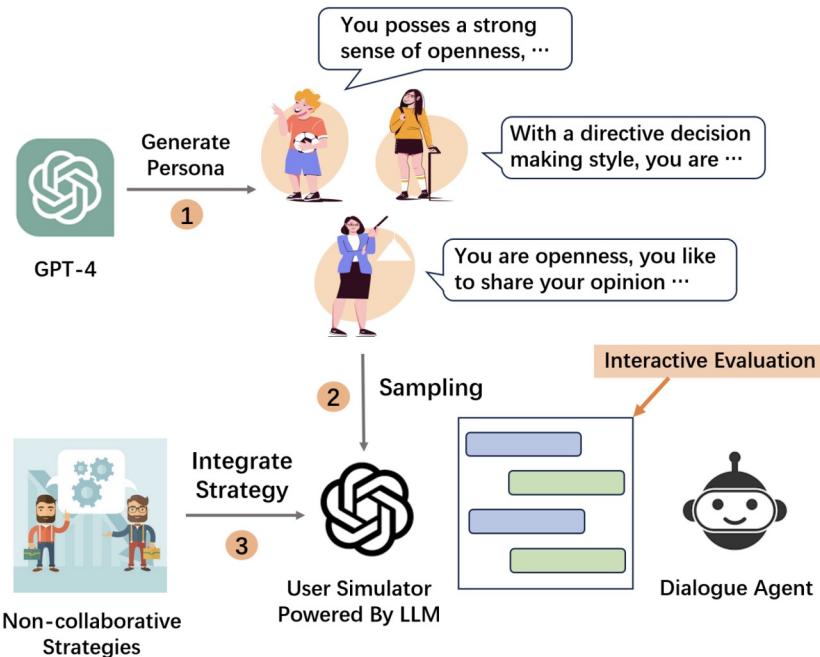


Why do we need to simulate diverse users?

Examples: Non-collaborative Dialogues (Negotiation/Persuasion)

- ❑ Existing dialogue systems overlook the integration of explicit **user-specific characteristics** in their strategic planning
- ❑ The training paradigm with a static user simulator fails to make strategic plans that can be **generalized to diverse users**

Role-playing Agents for Simulating Diverse Users



□ Big-Five Personality:

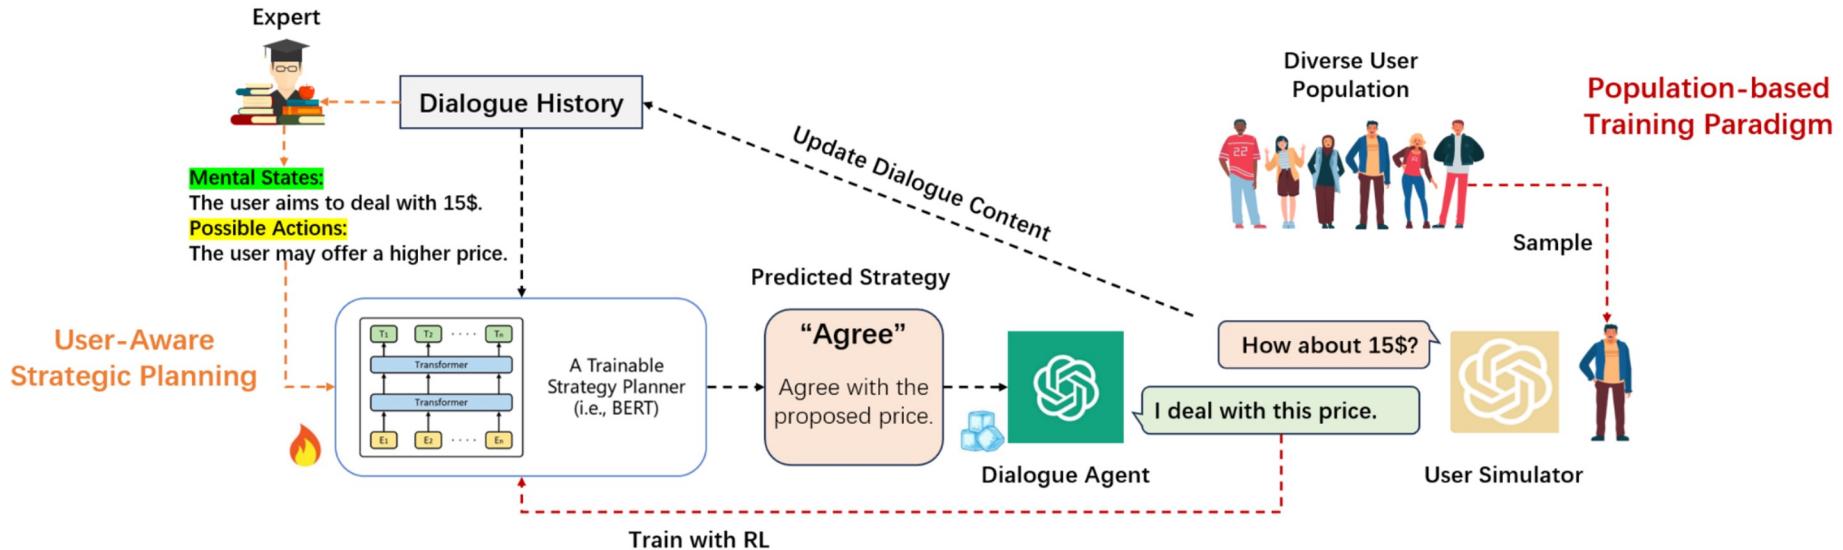
- *Openness, Conscientiousness, Extraversion, Agreeableness, and Neuroticism*

□ Decision-Making Styles:

- *Directive, Conceptual, Analytical, and Behavioral.*

	Personas	Price Negotiation			Persuasion for Good	
		SR↑	AT↓	SL%↑	SR↑	AT↓
Big Five	Openness	0.76↑ _{0.23}	6.66↑ _{0.63}	0.34↑ _{0.12}	0.47↑ _{0.34}	8.92↑ _{1.00}
	Conscientiousness	0.69↑ _{0.25}	7.20↑ _{1.04}	0.27↑ _{0.06}	0.39↑ _{0.33}	8.90↑ _{1.10}
	Extraversion	0.74↑ _{0.16}	6.17↑ _{1.47}	0.39↑ _{0.15}	0.45↑ _{0.35}	8.73↑ _{1.25}
	Agreeableness	0.40↑ _{0.01*}	6.82↑ _{0.71}	0.28↑ _{0.06}	0.18↑ _{0.12}	9.85↑ _{0.13*}
	Neuroticism	0.31↓ _{0.02*}	6.81↑ _{1.12}	0.20↓ _{0.02*}	0.12↑ _{0.02*}	9.78↑ _{1.44*}
Decision	Analytical	0.37↑ _{0.04*}	7.07↑ _{0.61}	0.20↑ _{0.06*}	0.16↑ _{0.09}	9.43↑ _{0.56*}
	Directive	0.41↑ _{0.05*}	6.71↑ _{1.48}	0.18↓ _{0.03*}	0.12↓ _{0.02*}	9.31↑ _{0.62}
	Behavioral	0.78↑ _{0.25}	6.45↑ _{1.20}	0.39↑ _{0.16}	0.53↑ _{0.37}	8.94↑ _{1.04}
	Conceptual	0.77↑ _{0.23}	6.62↑ _{0.78}	0.42↑ _{0.17}	0.49↑ _{0.36}	9.02↑ _{0.94}
Overall Performance		0.58↑ _{0.14}	6.72↑ _{1.01}	0.31↑ _{0.09}	0.32↑ _{0.23}	9.20↑ _{0.76}

Role-playing Agents for Simulating Diverse Users



New Training Paradigm with Diverse Simulated Users

- ❑ **User-aware Strategy Planning:** Predict user mental states and possible actions
- ❑ **Population-based Reinforcement Learning:** Sample a diverse group of simulated users to interact

Role-playing Agents for Simulating Diverse Users



Besides model learning, how about evaluation with simulated diverse users?

Wang et al., (2023) conclude that LLM-based user simulators are easier to accept the recommended items than human users during the evaluation of conversational recommender systems, since LLMs tend to follow the given instructions. → **Biased Evaluation!!!**

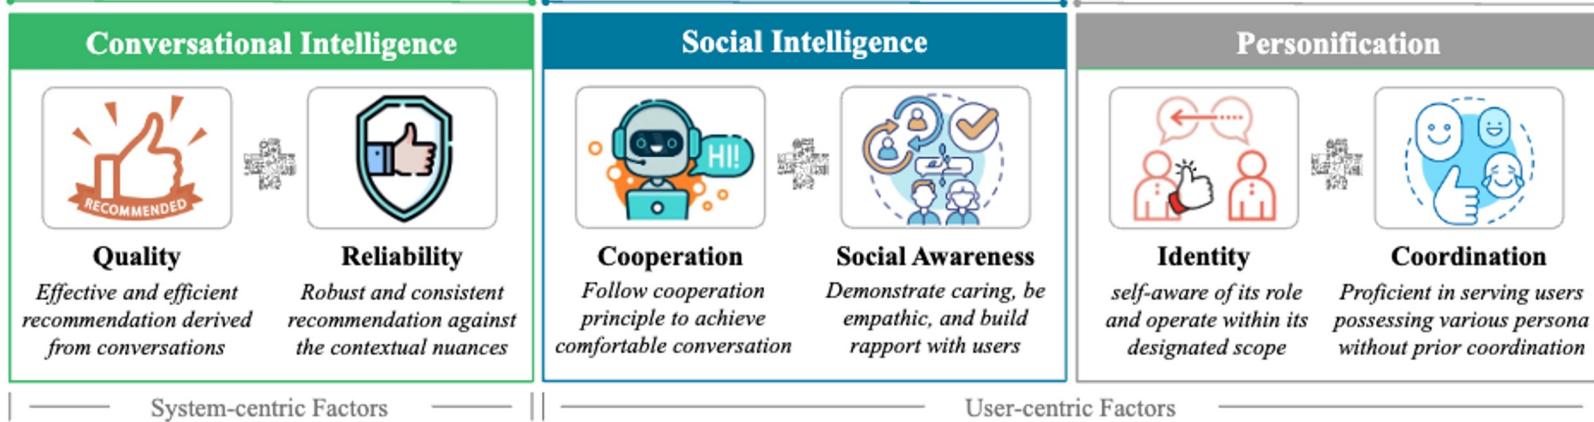
Persona	Templates (The Input of ChatGPT Paraphraser)	ChatGPT-paraphrased Persona Descriptions
Emotion=Boredom Age group=Adults	you are a person that are easy to be Boredom. This means that your are Feeling uninterested or uninspired by the recommended movie choices. Also, you are a Adults person	You are easily bored, feeling uninterested or uninspired by the recommended movie choices. As an adult, you seek movies that can captivate your attention.
Emotion=Anticipation Age group=Children	you are a person that are easy to be Anticipation. This means that your are Looking forward to watching recommended movies and experiencing new stories. Also, you are a Children person	You are filled with anticipation, looking forward to watching recommended movies and experiencing new stories. As a child, you enjoy the excitement of discovering new films.

Role-playing Agents for Simulating Diverse Users

Learn from conversations and evolve toward making recommendations as the conversation advances

Produce adequate social behavior for the recommendation during the conversation

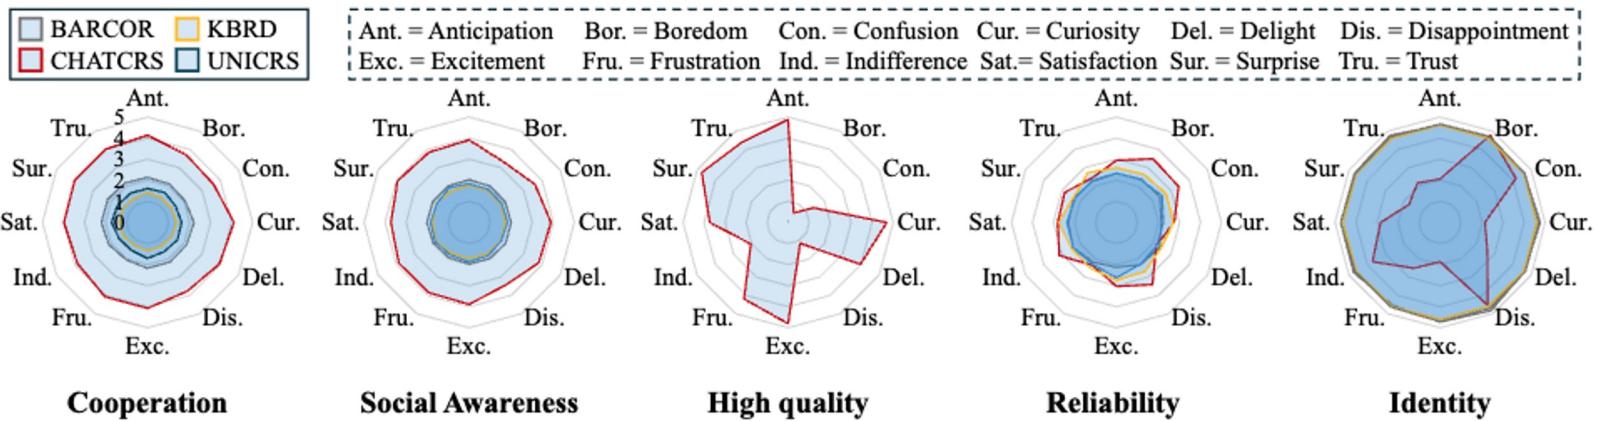
Perceive the identity of itself and the personality representation of users



Coordination

- **Definition:** Proficient in serving various and unknown users without prior coordination.
- **Metrics:** Computational metrics using the range and mean of other ability-specific scores that are calculated among various users.

Role-playing Agents for Simulating Diverse Users



Evaluation with Simulated Users from Different Personas

- Most CRS models, except for CHATCRS, show poor performance in sensing the variation of users.
- CHATCRS can properly deal with users' negative emotions, such as bored, confused, or disappointed.
- CHATCRS adopts sales pitches with deceptive tactics to persuade optimistic users to accept recommendations (Identity).

Overview of LLM-powered Conversational Agents



Profile

LLM-powered Conversational Agents for **User Simulation**



Memory

LLM-powered Conversational Agents for **Long-context Dialogues**



Planning

LLM-powered Conversational Agents for **Proactive Dialogues**



Action

LLM-powered Conversational Agents for **Real-world Problem Solving**

What is Long-context Dialogue?

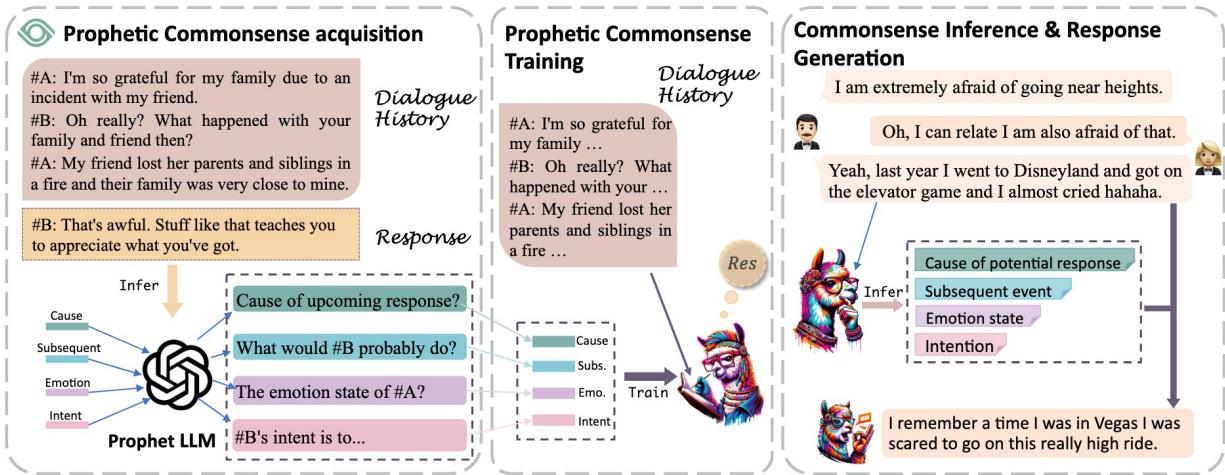


- Existing dialogue systems often concentrate on **single-session** interactions, overlooking the need for continuity in real-world conversational environments.

- Long-context dialogue systems requires memorization and personalization in **multi-session** conversations, providing more consistent and tailored responses.

External Knowledge for Long-context Dialogue

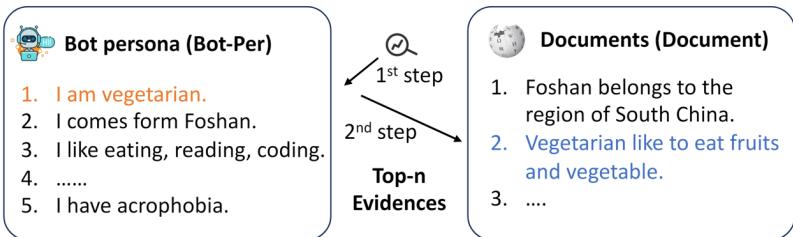
External Knowledge can act as supplementary guidance for the reasoning process.



The framework of employing external knowledge to reasoning.

Knowledge Sources:

- ❑ Commonsense Knowledge
- ❑ Medical Knowledge
- ❑ Psychology Knowledge
- ❑ ...



Wang et al., 2023. "Enhancing empathetic and emotion support dialogue generation with prophetic commonsense inference"

Wang et al., 2024. "UniMS-RAG: A Unified Multi-source Retrieval-Augmented Generation for Personalized Dialogue Systems"

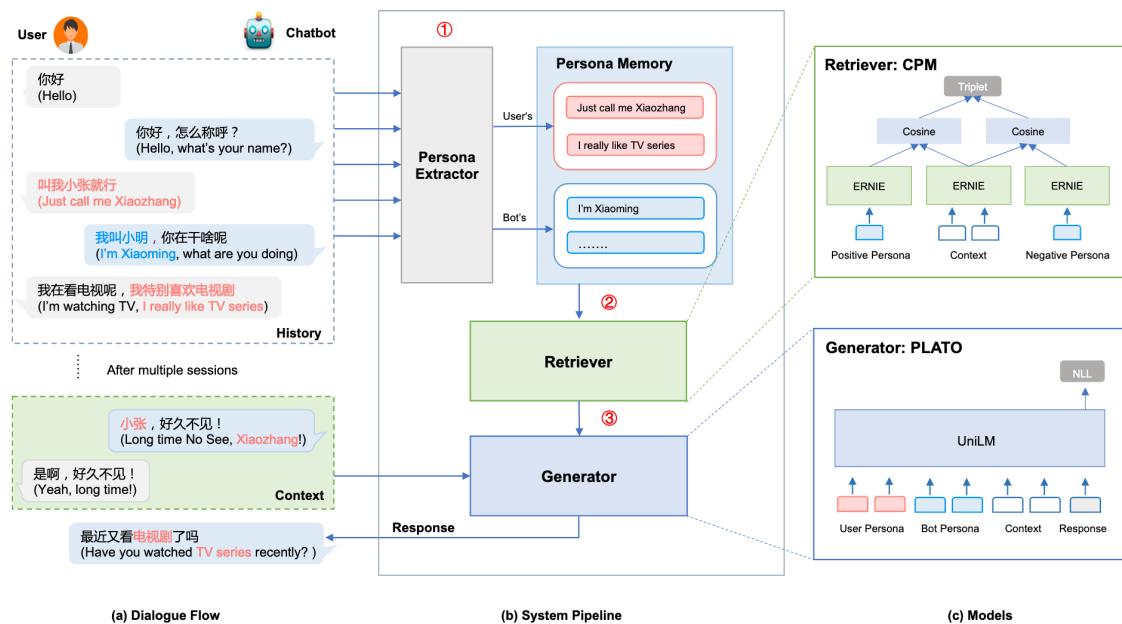
Internal Knowledge for Long-context Dialogue

- * Personas & Historical Events

Personas ensure the character consistency in long-context conversations.

Common Paradigm:

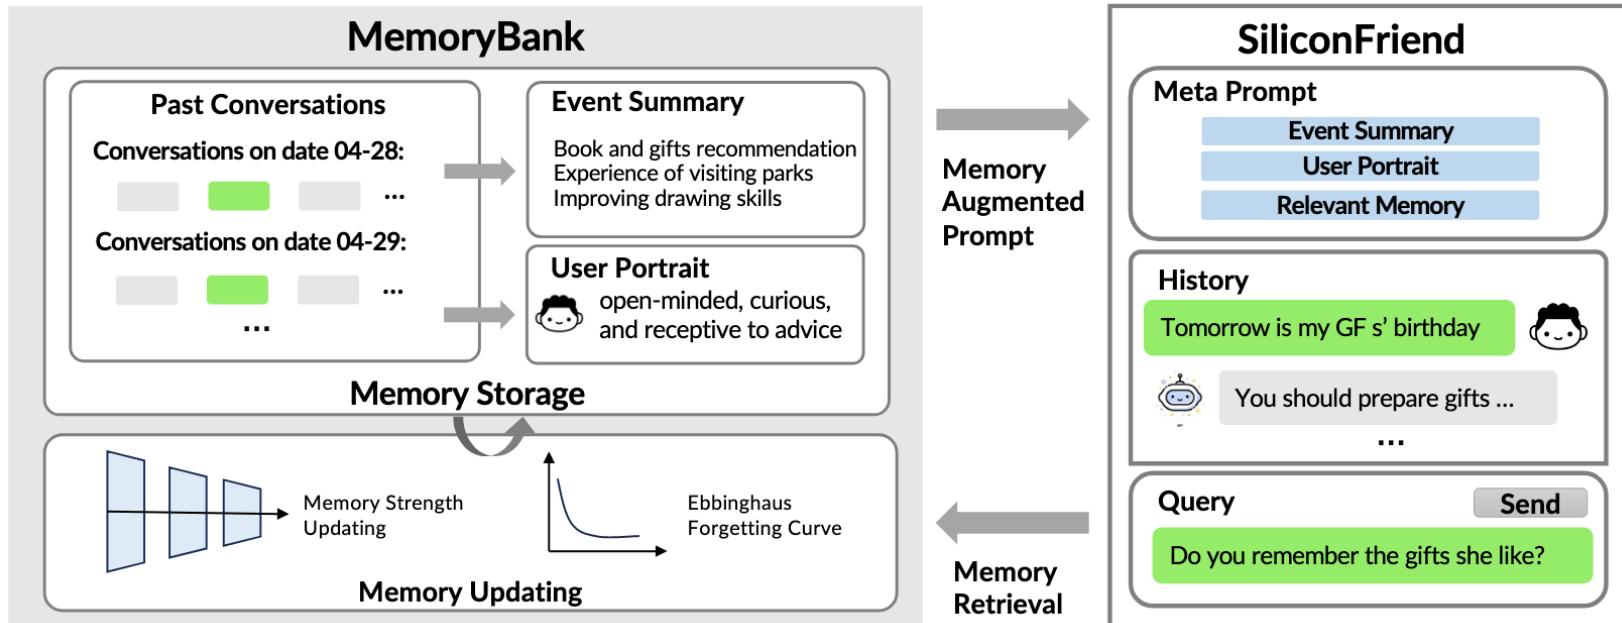
Typically, a **persona extraction** module is used to continuously **update persona** memory banks for both the user and the agent.



Internal Knowledge for Long-context Dialogue

- * Personas & Historical Events

Historical Events ensures dialogue coherence across sessions in long-context conversations.



Overview of LLM-powered Conversational Agents



Profile

LLM-powered Conversational Agents for **User Simulation**



Memory

LLM-powered Conversational Agents for **Long-context Dialogues**



Planning

LLM-powered Conversational Agents for **Proactive Dialogues**



Action

LLM-powered Conversational Agents for **Real-world Problem Solving**

Limitations of LLM-based Conversational Systems



Research ▾ API ▾ ChatGPT ▾ Safety Company ▾

Limitations

- ChatGPT sometimes writes plausible-sounding but incorrect or nonsensical answers. Fixing this issue is challenging, as: (1) during RL training, there's currently no source of truth; (2) training the model to be more cautious causes it to decline questions that it can answer correctly; and (3) supervised training misleads the model because the ideal answer depends on what the model knows, rather than what the human demonstrator knows.
- ChatGPT is sensitive to tweaks to the input phrasing or attempting the same prompt multiple times. For example, given one phrasing of a question, the model can claim to not know the answer, but given a slight rephrase, can answer correctly.
- The model is often excessively verbose and overuses certain phrases, such as restating that it's a language model trained by OpenAI. These issues arise from biases in the training data (trainers prefer longer answers that look more comprehensive) and well-known over-optimization issues.^{1, 2}
- Ideally, the model would ask clarifying questions when the user provided an ambiguous query. Instead, our current models usually guess what the user intended.
- While we've made efforts to make the model refuse inappropriate requests, it will sometimes respond to harmful instructions or exhibit biased behavior.

Limitations of LLM-based Conversational Systems



Research ▾ API ▾ ChatGPT ▾ Safety Company ▾

Limitations

- ChatGPT sometimes writes plausible-sounding but incorrect or nonsensical answers. Fixing this issue is challenging, as: (1) during PL training, there's currently

- Ideally, the model would ask clarifying questions when the user provided an ambiguous query. Instead, our current models usually guess what the user intended.
- While we've made efforts to make the model refuse inappropriate requests, it will sometimes respond to harmful instructions or exhibit biased behavior.

biases in the training data (trainers prefer longer answers that look more comprehensive) and well-known over-optimization issues.^{1, 2}

- Ideally, the model would ask clarifying questions when the user provided an ambiguous query. Instead, our current models usually guess what the user intended.
- While we've made efforts to make the model refuse inappropriate requests, it will sometimes respond to harmful instructions or exhibit biased behavior.

- ★ **Instruction-following/Reactive** Conversational AI – The conversation is led by the user, and the system simply follows the user's instructions or intents.

Proactive Conversational Agent

A proactive conversational agent is a conversational system that can **plan** the conversation to achieve the conversational goals by taking **initiative** and **anticipating** long-term impacts on themselves or human users.

Goal Awareness for Conversational AI: Proactivity, Non-collaborativity, and Beyond

Yang Deng, Wenqiang Lei, Minlie Huang, Tat-Seng Chua

ACL 2023 Tutorial



Anticipation

To anticipate future impacts on the task or human users.

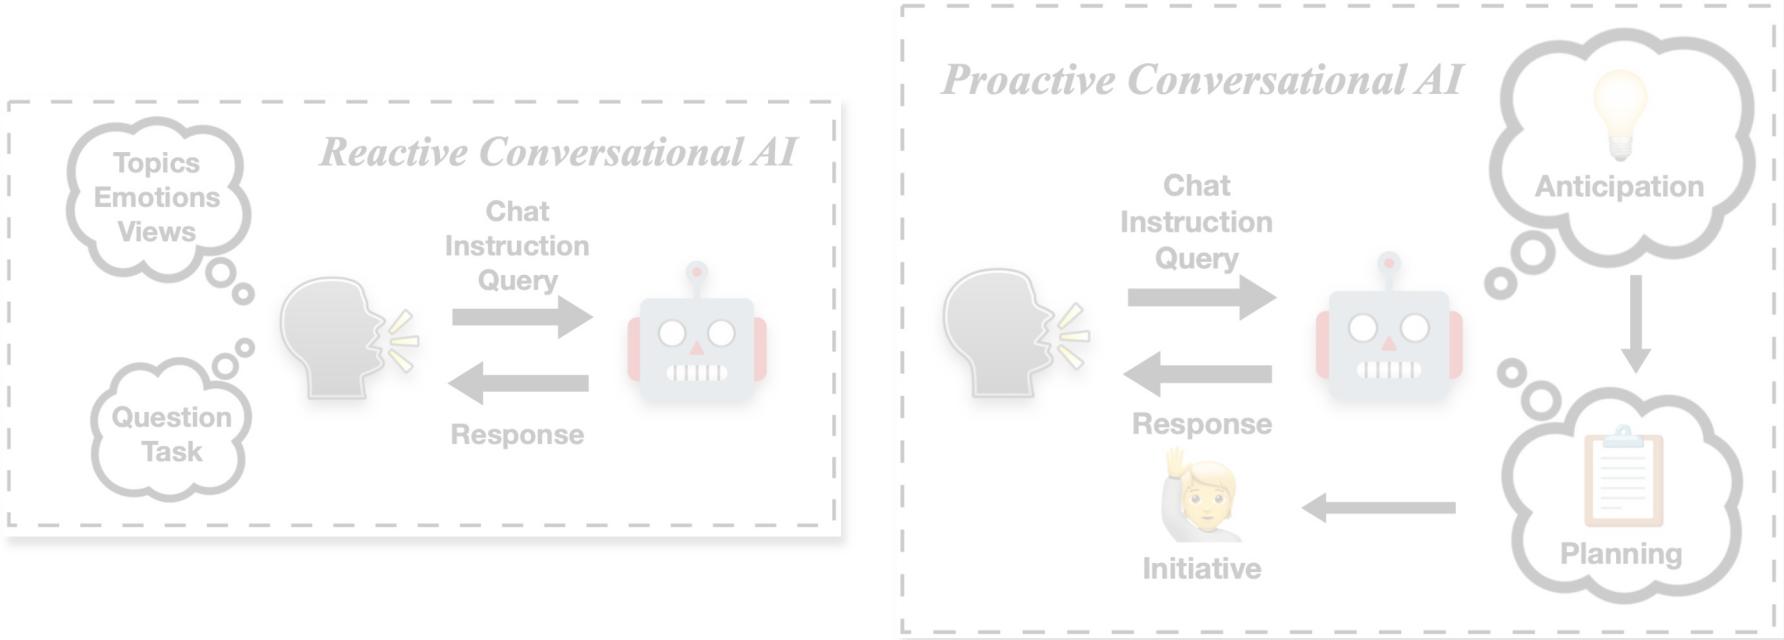
Initiative

To take fine-grained and diverse initiative behaviours.

Planning

To effectively and efficiently guide the conversation towards the goal.

Reactive vs. Proactive Conversational AI



Triggering the Proactivity of LLMs via In-Context Learning



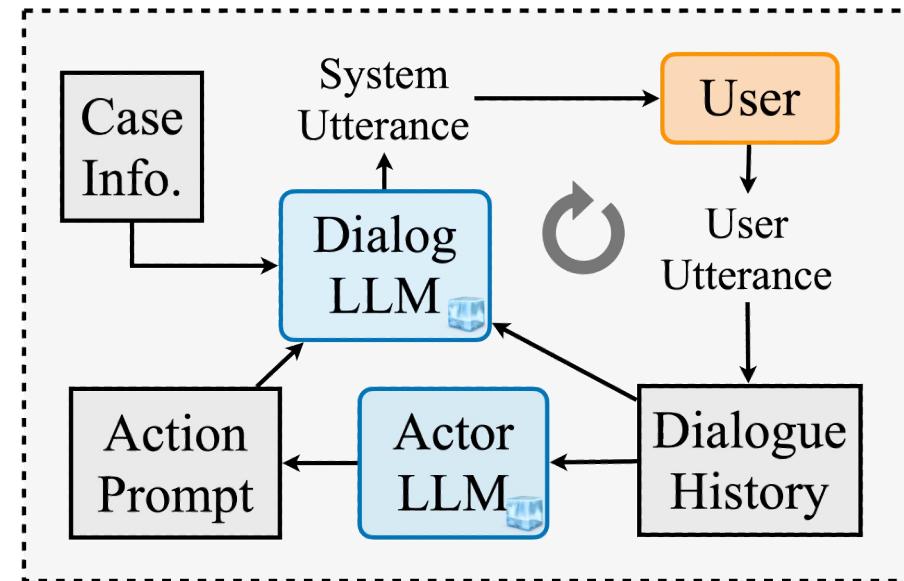
Can LLM-based Conversational Agents effectively handle proactive dialogue problems without fine-tuning?

❑ Advantages of In-Context Learning

- ✓ Training-free
- ✓ Easy-to-apply

➤ Proactive Chain-of-Thought

- * Fine-grained Initiative
- * Intermediate Reasoning



Proactive Chain-of-Thought Prompting (ProCoT)

□ Standard Prompting

- Input: Task Background & Conversation History
- Output: Response

$$p(r|\mathcal{D}, \mathcal{C})$$

(1) Clarification Dialogues: Abg-CoQA

Task Background: The grounded document is "Angie She made a drawing of her mother. Her mother found a large red book. Then they went to the Mystery section. Angie sat in a blue chair. She drew a picture of her brother. Her mother found the book. It was a green book. ..."

Conversation History: [{"User": "What did she draw?", "System": "Her mother", "User": "What did her mother find?", "System": "The book", "User": "What color was it?"}]

(1a) Standard

Prompt: Given the task background and the conversation history, please generate the response:
Response: Green X

Proactive Chain-of-Thought Prompting (ProCoT)

Standard Prompting

- Input: Task Background & Conversation History
- Output: Response

$$p(r|\mathcal{D}, \mathcal{C})$$

Proactive Prompting

- Input: + Action Space
- Output: + Action

$$p(a, r|\mathcal{D}, \mathcal{C}, \mathcal{A})$$

(1) Clarification Dialogues: Abg-CoQA

Task Background: The grounded document is "Angie She made a drawing of her mother. Her mother found a large red book. Then they went to the Mystery section. Angie sat in a blue chair. She drew a picture of her brother. Her mother found the book. It was a green book."

Conversation History: ["User": "What did she draw?", "System": "Her mother", "User": "What did her mother find?", "System": "The book", "User": "What color was it?"]

(1a) Standard

Prompt: Given the task background and the conversation history, please generate the response:
Response: Green X

(1b) Proactive

Act: ["Directly Answer", "Ask a Clarification Question"]
Prompt: Given the task background and the conversation history, please use appropriate actions to generate the response:
Response: Ask a clarification question: Could you provide more information? X

Proactive Chain-of-Thought Prompting (ProCoT)

Standard Prompting

- Input: Task Background & Conversation History
- Output: Response

$$p(r|\mathcal{D}, \mathcal{C})$$

Proactive Prompting

- Input: + Action Space
- Output: + Action

$$p(a, r|\mathcal{D}, \mathcal{C}, \mathcal{A})$$

Proactive Chain-of-Thought Prompting

- Output: + Reasoning Chain

$$p(t, a, r|\mathcal{D}, \mathcal{C}, \mathcal{A})$$

(1) Clarification Dialogues: Abg-CoQA

Task Background: The grounded document is "Angie She made a drawing of her mother. Her mother found a large red book. Then they went to the Mystery section. Angie sat in a blue chair. She drew a picture of her brother. Her mother found the book. It was a green book."

Conversation History: ["User": "What did she draw?", "System": "Her mother", "User": "What did her mother find?", "System": "The book", "User": "What color was it?"]

(1a) Standard

Prompt: Given the task background and the conversation history, please generate the response:
Response: Green X

(1c) Proactive CoT

Act: ["Directly Answer", "Ask a Clarification Question"]

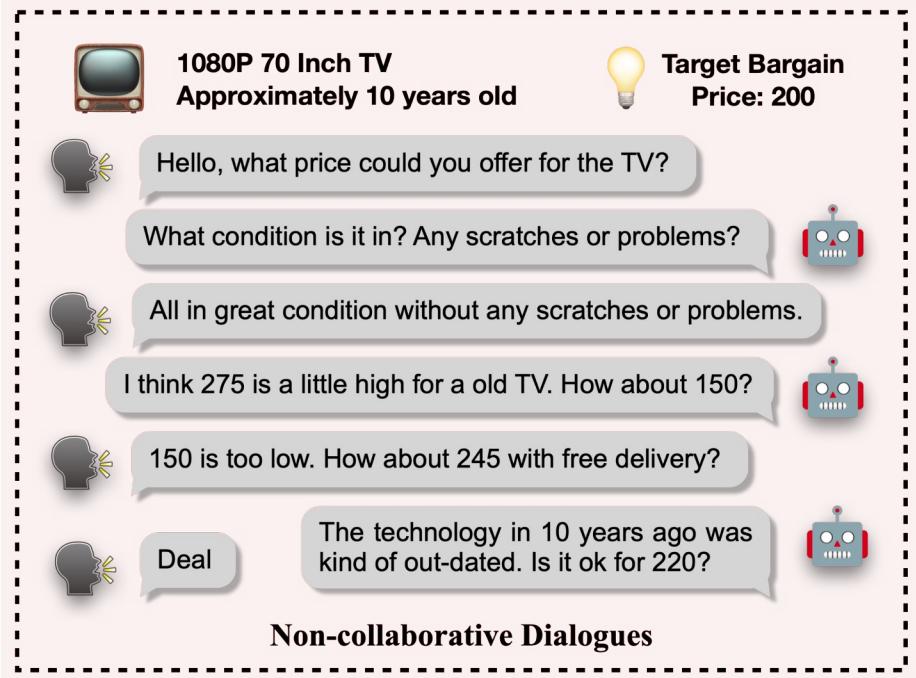
Prompt: Given the task background and the conversation history, first analyse whether the question at the current turn is ambiguous to answer, and then select appropriate actions to generate the response:

Response: There are two books that Angie's mother found. It is unclear which book is referred to. Therefore, the question is ambiguous. Ask a clarification question: Which book are you refer to? ✓

(1b) Proactive

Act: ["Directly Answer", "Ask a Clarification Question"]
Prompt: Given the task background and the conversation history, please use appropriate actions to generate the response:
Response: Ask a clarification question: Could you provide more information? X

Evaluating LLMs on three Proactive Dialogue Problems



Evaluation of Clarification in Information-seeking Dialogues

Method	Shot	Prompt	Open-domain			Finance		
			Abg-CoQA			PACIFIC		
			CNP	CQG	CNP	CQG		
Baseline	-	-	22.1	36.5	30.0	79.0	69.2	38.2
SOTA	-	-	23.6	38.2	56.0	86.9	90.7	80.1
Vicuna-13B	0	Standard	-	11.3	0.0	-	1.2	0.0
	1	Standard	-	11.4	0.0	-	2.5	0.0
	0	Proactive	4.1	13.2	0.0	2.3	2.3	0.0
	1	Proactive	12.1	13.2	4.5	0.0	3.3	0.0
	0	ProCoT	1.4	21.3	9.1	9.7	3.8	10.5
ChatGPT	1	ProCoT	18.3	23.7	22.7	27.0	41.3	33.1
	0	Standard	-	12.1	0.0	-	2.2	0.0
	1	Standard	-	12.3	0.0	-	2.0	0.0
	0	Proactive	22.0	13.7	17.6	19.4	2.9	0.0
	1	Proactive	20.4	23.4	23.5	17.7	14.0	12.5
	0	ProCoT	23.8	21.6	32.4	28.0	21.5	26.7
	1	ProCoT	27.9	18.4	45.9	27.7	16.2	35.8



LLMs barely ask clarification questions.

Evaluation of Clarification in Information-seeking Dialogues

Method	Shot	Prompt	Open-domain			Finance		
			Abg-CoQA			PACIFIC		
			CNP	CQG	CNP	CQG		
Baseline	-	-	22.1	36.5	30.0	79.0	69.2	38.2
SOTA	-	-	23.6	38.2	56.0	86.9	90.7	80.1
Vicuna-13B	0	Standard	-	11.3	0.0	-	1.2	0.0
	1	Standard	-	11.4	0.0	-	2.5	0.0
	0	Proactive	4.1	13.2	0.0	2.3	2.3	0.0
	1	Proactive	12.1	13.2	4.5	0.0	3.3	0.0
	0	ProCoT	1.4	21.3	9.1	9.7	3.8	10.5
	1	ProCoT	18.3	23.7	22.7	27.0	41.3	33.1
ChatGPT	0	Standard	-	12.1	0.0	-	2.2	0.0
	1	Standard	-	12.3	0.0	-	2.0	0.0
	0	Proactive	22.0	13.7	17.6	19.4	2.9	0.0
	1	Proactive	20.4	23.4	23.5	17.7	14.0	12.5
	0	ProCoT	23.8	21.6	32.4	28.0	21.5	26.7
	1	ProCoT	27.9	18.4	45.9	27.7	16.2	35.8



LLMs barely ask clarification questions.



ProCoT largely overcomes this issue in open-domain, but the performance is still unsatisfactory in domain-specific applications.

Evaluation on Target-guided Chit-chat Dialogues

Method	Shot	Prompt	Easy Target			Hard Target		
			Succ.(%)	Turns	Coh.	Succ.(%)	Turns	Coh.
GPT2	-	-	22.3	2.86	0.23	17.3	2.94	0.21
DKRN	-	-	38.6	4.24	0.33	21.7	7.19	0.31
CKC	-	-	41.9	4.08	0.35	24.8	6.88	0.33
TopKG	-	-	48.9	3.95	0.31	27.3	4.96	0.33
COLOR	-	-	66.3	-	0.36	30.1	-	0.35
Vicuna-13B	0	Standard	63.0	2.63	0.43	62.5	2.45	0.39
	1	Standard	62.7	2.83	0.45	65.0	2.90	0.43
	0	Proactive	37.8	2.71	0.48	35.6	2.56	0.55
	1	Proactive	48.3	2.71	0.50	34.6	2.95	0.51
	0	ProCoT	65.2	4.22	0.49	54.9	4.17	0.45
	1	ProCoT	72.3	3.55	0.52	59.8	3.81	0.48
ChatGPT	0	Standard	97.5	2.26	0.38	96.3	2.30	0.41
	1	Standard	96.3	2.42	0.42	93.5	2.28	0.38
	0	Proactive	85.9	3.20	0.47	83.0	2.83	0.43
	1	Proactive	90.7	2.86	0.36	86.2	2.94	0.31
	0	ProCoT	96.3	2.47	0.41	92.0	2.29	0.34
	1	ProCoT	95.9	2.63	0.45	92.1	2.47	0.39



LLMs are proficient at performing topic shifting towards the designated target.

Evaluation on Target-guided Chit-chat Dialogues

Method	Shot	Prompt	Easy Target			Hard Target		
			Succ.(%)	Turns	Coh.	Succ.(%)	Turns	Coh.
GPT2	-	-	22.3	2.86	0.23	17.3	2.94	0.21
DKRN	-	-	38.6	4.24	0.33	21.7	7.19	0.31
CKC	-	-	41.9	4.08	0.35	24.8	6.88	0.33
TopKG	-	-	48.9	3.95	0.31	27.3	4.96	0.33
COLOR	-	-	66.3	-	0.36	30.1	-	0.35
Vicuna-13B	0	Standard	63.0	2.63	0.43	62.5	2.45	0.39
	1	Standard	62.7	2.83	0.45	65.0	2.90	0.43
	0	Proactive	37.8	2.71	0.48	35.6	2.56	0.55
	1	Proactive	48.3	2.71	0.50	34.6	2.95	0.51
	0	ProCoT	65.2	4.22	0.49	54.9	4.17	0.45
	1	ProCoT	72.3	3.55	0.52	59.8	3.81	0.48
ChatGPT	0	Standard	97.5	2.26	0.38	96.3	2.30	0.41
	1	Standard	96.3	2.42	0.42	93.5	2.28	0.38
	0	Proactive	85.9	3.20	0.47	83.0	2.83	0.43
	1	Proactive	90.7	2.86	0.36	86.2	2.94	0.31
	0	ProCoT	96.3	2.47	0.41	92.0	2.29	0.34
	1	ProCoT	95.9	2.63	0.45	92.1	2.47	0.39

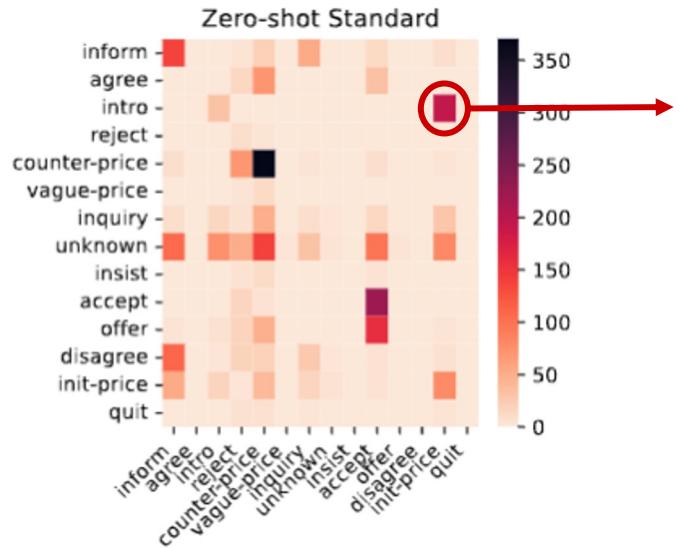


LLMs are proficient at performing topic shifting towards the designated target.



LLMs tend to make aggressive topic transition.

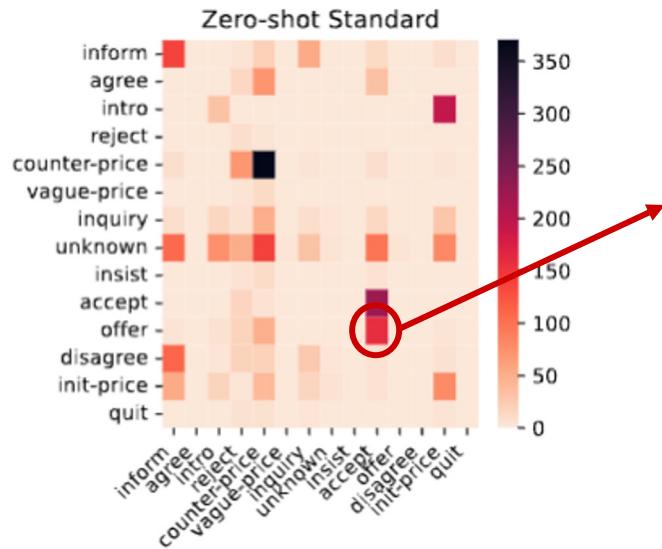
Evaluation on Non-collaborative Dialogues (Negotiation)



- Tends to propose the initial price (**init-price**) instead of greetings (**intro**) at the beginning.
- Often directly accepts the buyer's offer (**accept**) when it is supposed to offer another price for negotiation (**offer**).
- Tends to propose a counter price (**counter-price**) to make compromise with the user.

Relationships between reference and predicted negotiation strategies.

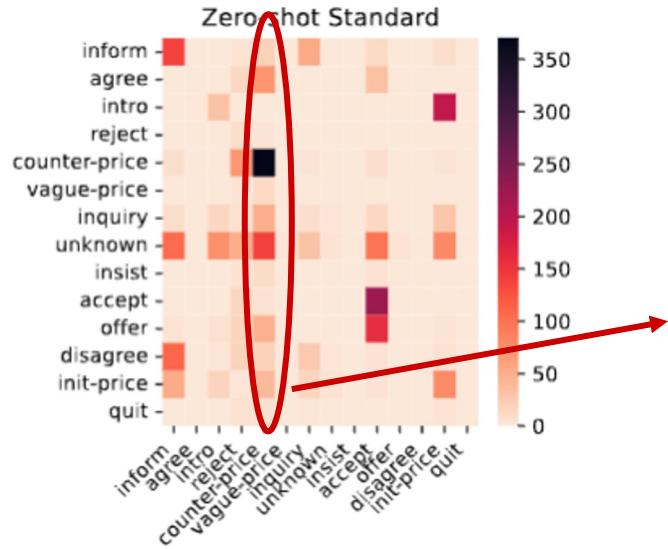
Evaluation on Non-collaborative Dialogues (Negotiation)



- Tends to propose the initial price (**init-price**) instead of greetings (**intro**) at the beginning.
- Often directly accepts the buyer's offer (**accept**) when it is supposed to offer another price for negotiation (**offer**).
- Tends to propose a counter price (**counter-price**) to make compromise with the user.

Relationships between reference and predicted negotiation strategies.

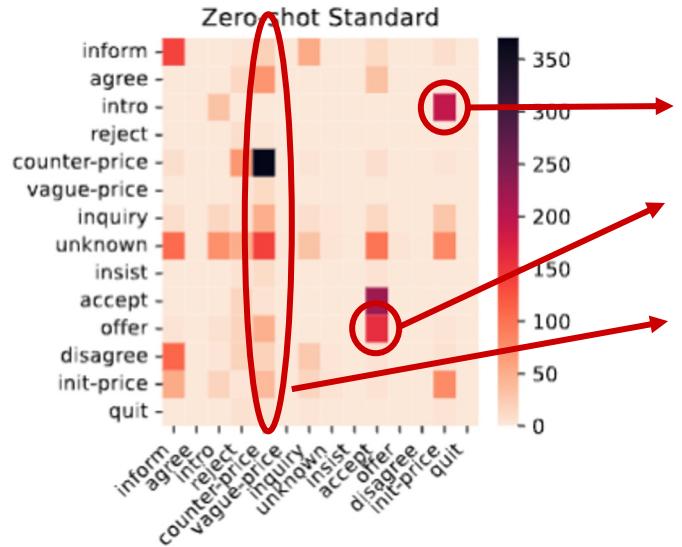
Evaluation on Non-collaborative Dialogues (Negotiation)



- Tends to propose the initial price (**init-price**) instead of greetings (**intro**) at the beginning.
- Often directly accepts the buyer's offer (**accept**) when it is supposed to offer another price for negotiation (**offer**).
- Tends to propose a counter price (**counter-price**) to make compromise with the user.

Relationships between reference and predicted negotiation strategies.

Evaluation on Non-collaborative Dialogues (Negotiation)



- Tends to propose the initial price (**init-price**) instead of greetings (**intro**) at the beginning.
- Often directly accepts the buyer's offer (**accept**) when it is supposed to offer another price for negotiation (**offer**).
- Tends to propose a counter price (**counter-price**) to make compromise with the user.



LLMs fail to make strategic decision for non-collaborative dialogues and tend to compromise with the user.

Relationships between reference and predicted negotiation strategies.

Lessons Learned from the Evaluation

❑ Clarification in Information-seeking Dialogue

- ❑ Barely ask clarification questions.
- ❑ Perform badly at domain-specific applications.

❑ Target-guided Open-domain Dialogue

- ❑ Proficient at topic shifting towards the designated target.
- ❑ Tend to make aggressive topic transition.

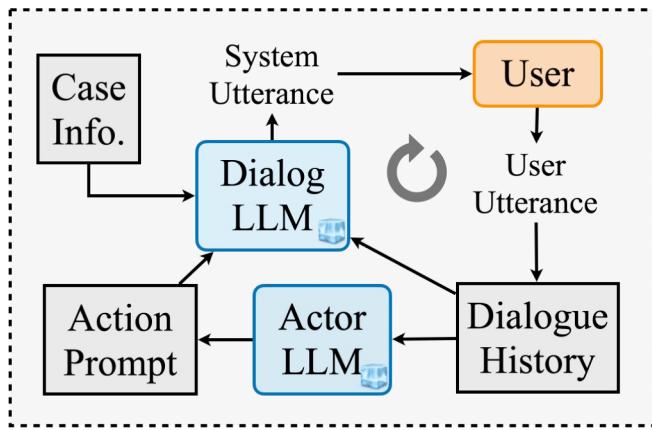
❑ Non-collaborative Dialogue

- ❑ Fail to make strategic plans.
- ❑ Tend to compromise with the user.



*LLM-based Conversational Agents
fail to plan appropriate initiative
behaviours.*

Limitations of In-context Learning Approaches



- ❑ Fail to optimize the long-term goal of the conversation.
- ❑ Not learnable.
- ❑ Limited by the strategy planning capability of LLMs.

➤ Reinforcement Learning with Goal-oriented AI Feedback

Problem Formulation

- Formulate the proactive conversation as a **Markov Decision Process (MDP)**.
- The objective is to learn a policy π maximizing the expected cumulative rewards over the observed dialogue episodes as:

$$\pi^* = \arg \max_{\pi \in \Pi} \left[\sum_{t=0}^T \mathcal{R}(s_t) \right]$$

Reward Function

$$= \arg \max_{\pi \in \Pi} \left[\sum_{t=0}^T \mathcal{R}(\mathcal{T}(s_{t-1}, a_t)) \right]$$

State Transition

$$= \arg \max_{\pi \in \Pi} \left[\sum_{t=0}^T \mathcal{R}(\mathcal{T}(s_{t-1}, \pi(s_{t-1}))) \right]$$

Policy Network



How to enable the policy learning with LLMs?

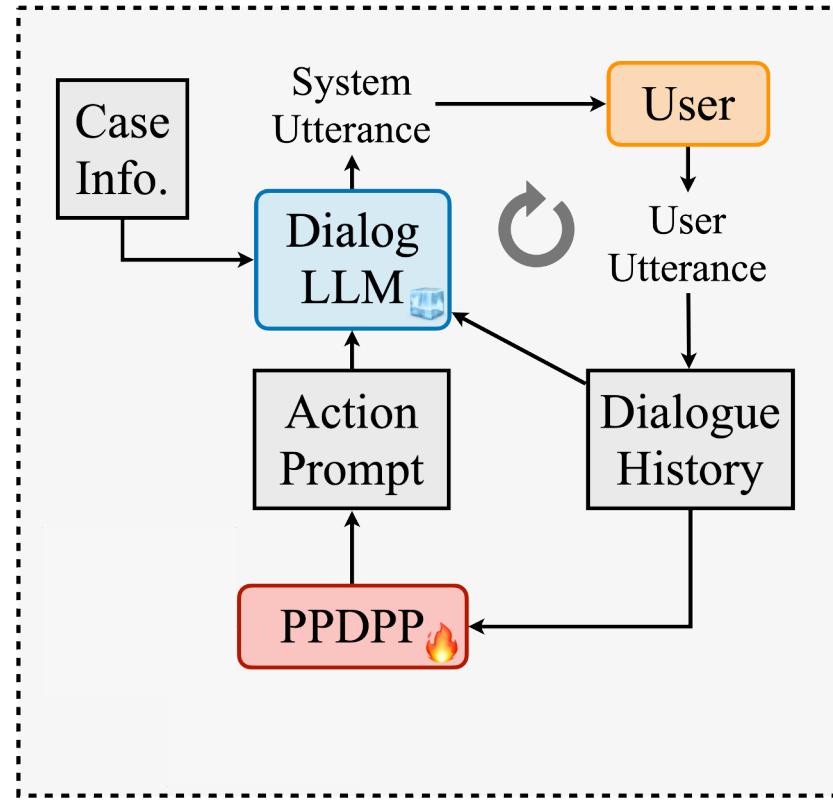
Policy Network – Plug-and-Play Dialogue Policy Planner

- A tunable language model plug-in for dialogue strategy learning.

$$a_t = \pi(s_{t-1})$$

- Conduct **Supervised Fine-Tuning** on available human-annotated corpus.

$$\mathcal{L}_c = -\frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} \frac{1}{T_d} \sum_{t=1}^{T_d} a_t \log y_t$$



Reward Function – Learning from AI Feedback

- An LLM as the reward model to assess the goal achievement and provide **goal-oriented AI feedback**.

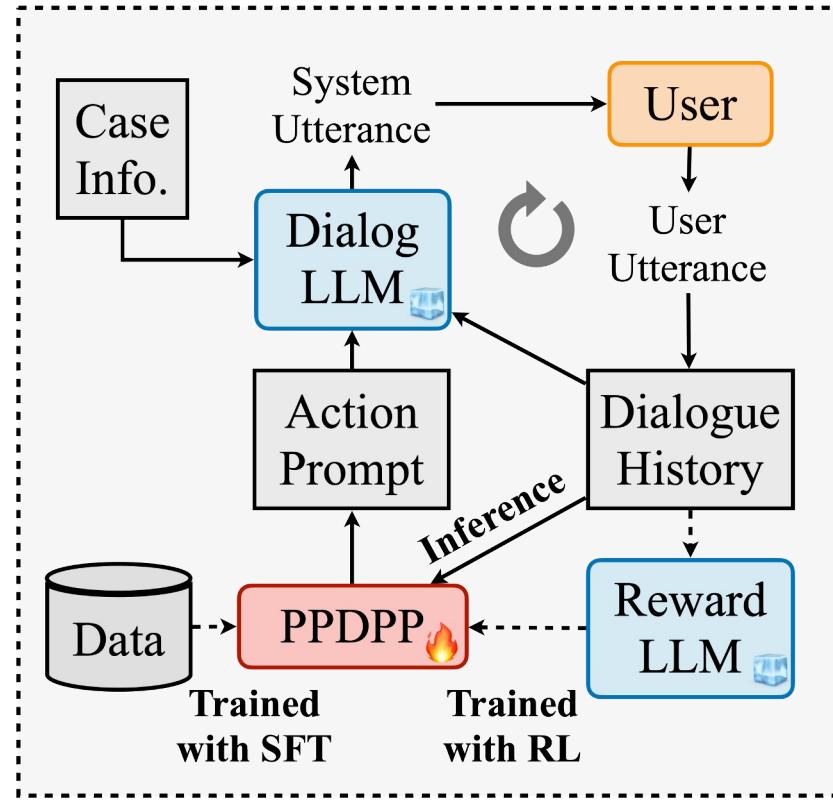
$$\mathcal{R}(s_t) = \frac{1}{l} \sum_{i=1}^l \mathcal{M}_r(\text{LLM}_{\text{rwd}}(p_{\text{rwd}}; s_t; \tau))$$

- Employ **Reinforcement Learning** to further tune the policy model.

$$\theta \leftarrow \theta - \alpha \nabla \log \pi_\theta(a_t | s_t) R_t$$



Interacting with real user is costly!



State Transition – Multi-agent Simulation

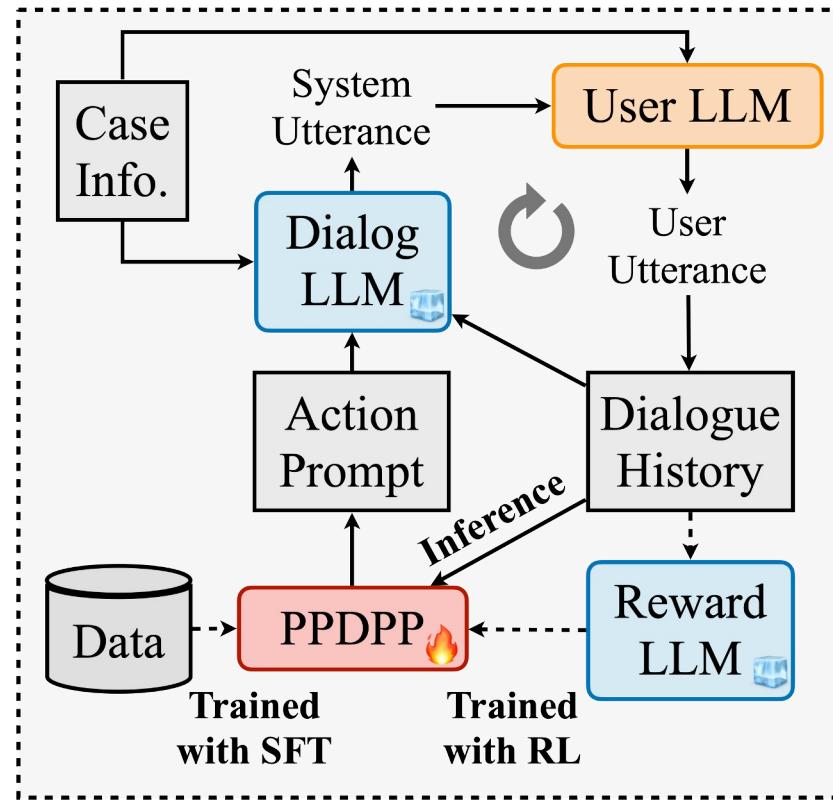
- ❑ An LLM to simulate the user with user profiles.
- ❑ Employ **Multi-agent Simulation** to collect dynamic interaction data.

$$u_t^{sys} = \text{LLM}_{sys}(p_{sys}; \mathcal{M}_a(a_t); s_{t-1})$$

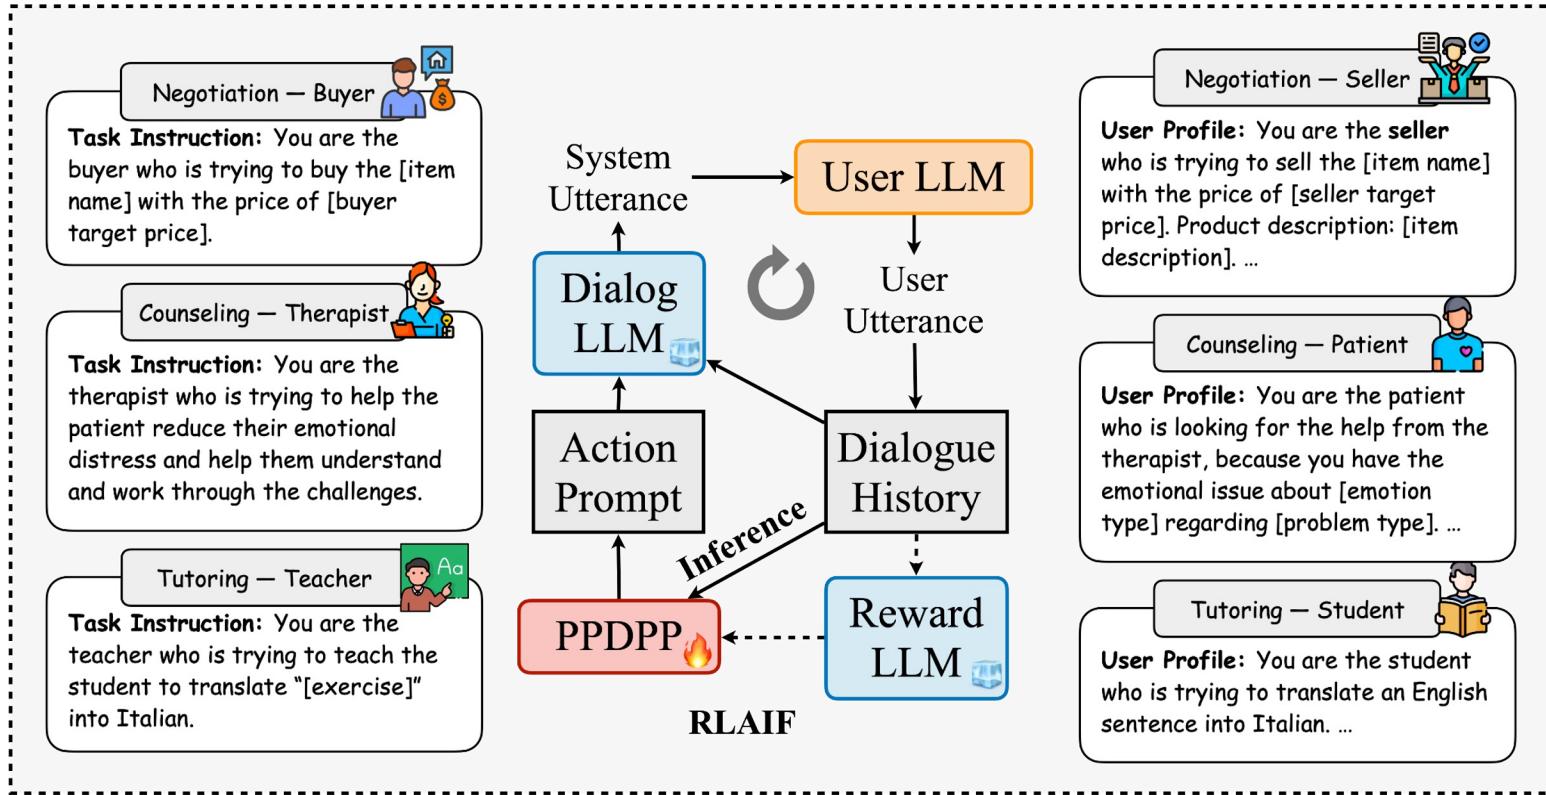
$$u_t^{usr} = \text{LLM}_{usr}(p_{usr}; s_{t-1}; u_t^{sys})$$

$$s_t = \mathcal{T}(s_{t-1}, a_t)$$

$$= \{s_{t-1}; u_t^{sys}, u_t^{usr}\}$$



Examples: Multi-agent Simulation



Overview of LLM-powered Conversational Agents



Profile

LLM-powered Conversational Agents for **User Simulation**



Memory

LLM-powered Conversational Agents for **Long-context Dialogues**



Planning

LLM-powered Conversational Agents for **Proactive Dialogues**



Action

LLM-powered Conversational Agents for **Real-world Problem Solving**

Web Agents

Web Agents aims to accomplish the tasks defined in natural language, such as booking tickets, through multi-step interactions with the web-grounded environment.

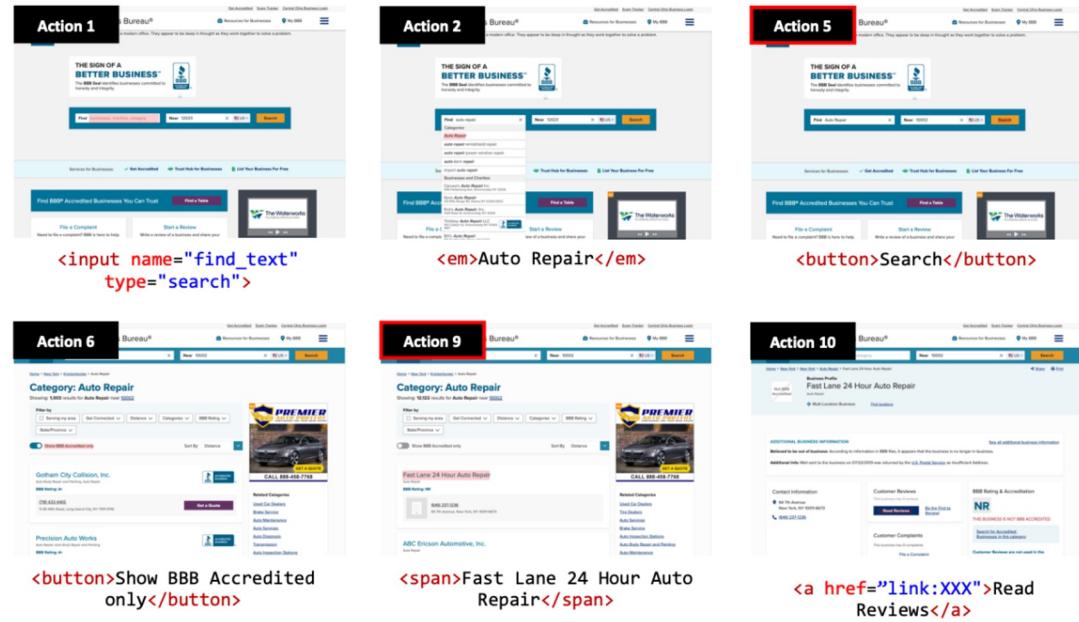
Task Description:

Show me the reviews for the auto repair business closest to 10002.

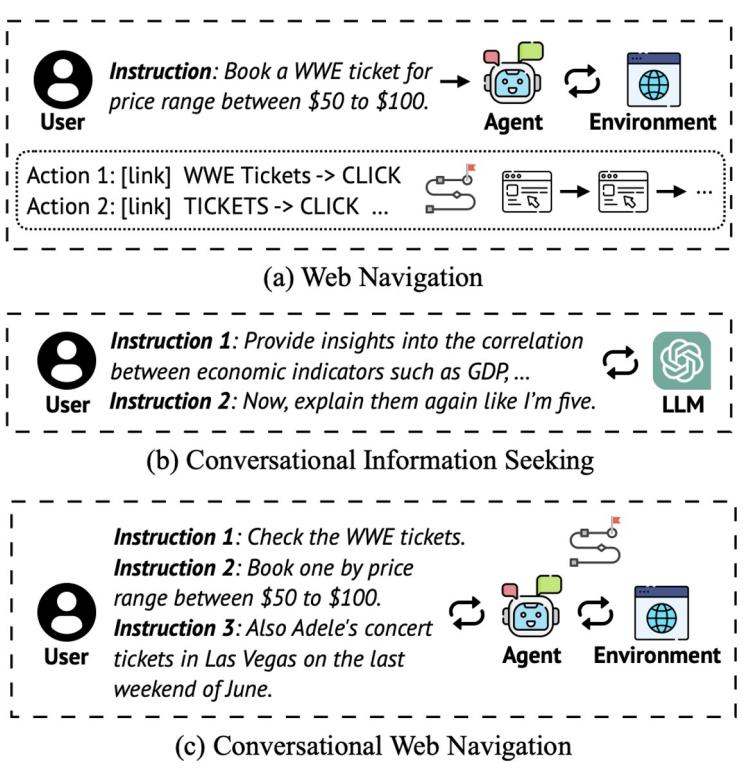
Action Sequence:

Target Element	Operation
1. [searchbox] <i>Find</i>	TYPE: <i>auto repair</i>
2. [button] <i>Auto Repair</i>	CLICK
3. [textbox] <i>Near</i>	TYPE: <i>10002</i>
4. [button] <i>10002</i>	CLICK
5. [button] <i>Search</i>	CLICK
6. [switch] <i>Show BBB Accredited only</i>	CLICK
7. [svg]	CLICK
8. [button] <i>Sort By</i>	CLICK
9. [link] <i>Fast Lane 24 Hour Auto Repair</i>	CLICK
10. [link] <i>Read Reviews</i>	CLICK

Webpage Snapshots:



Conversational Web Agents



Web Navigation

- Single-turn User Instruction
- Multi-step Environment Interaction

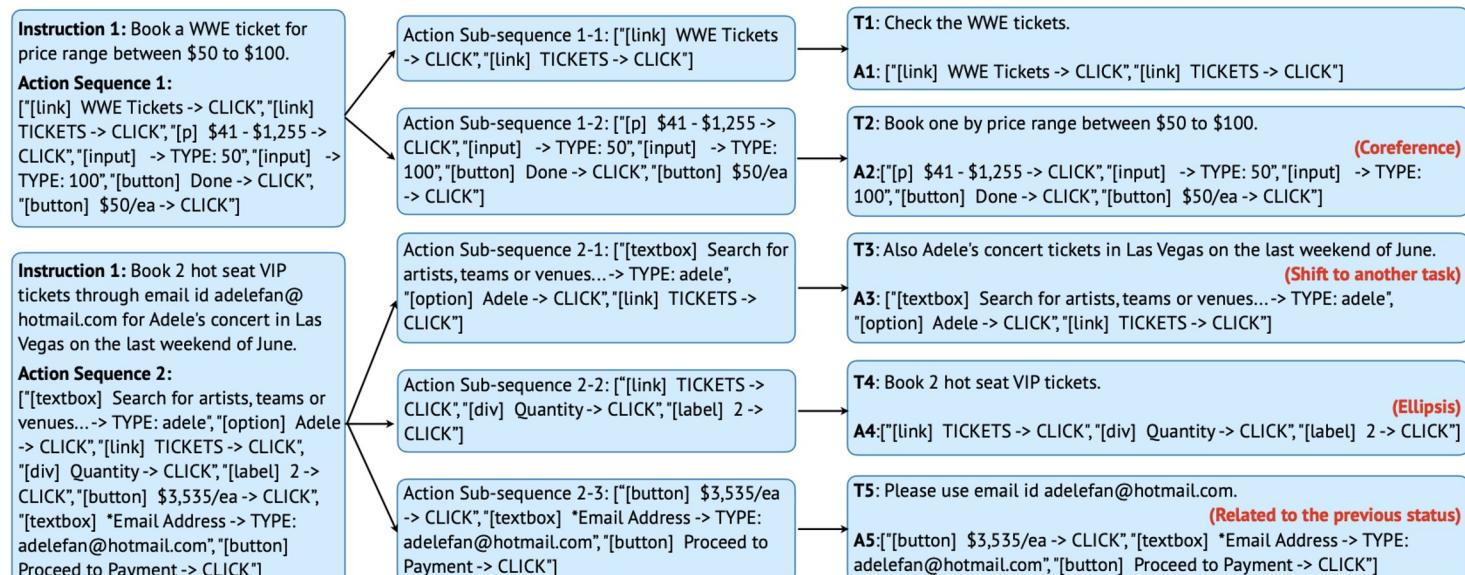
Conversational Information Seeking

- Multi-turn User Instruction
- No/Single-step Environment Interaction

Conversational Web Navigation

- Multi-turn User Instruction
- Multi-step Environment Interaction

Constructing the MT-Mind2Web Dataset



Organize Conversation Session

Decompose Complex Instructions

Rewrite Conversational Instructions



Organize Conversation Sessions

Decompose Complex Instructions

Rewrite Conversational Instructions

Modify



Challenges in Conversational Web Agents

<Longer and Noisier Context>

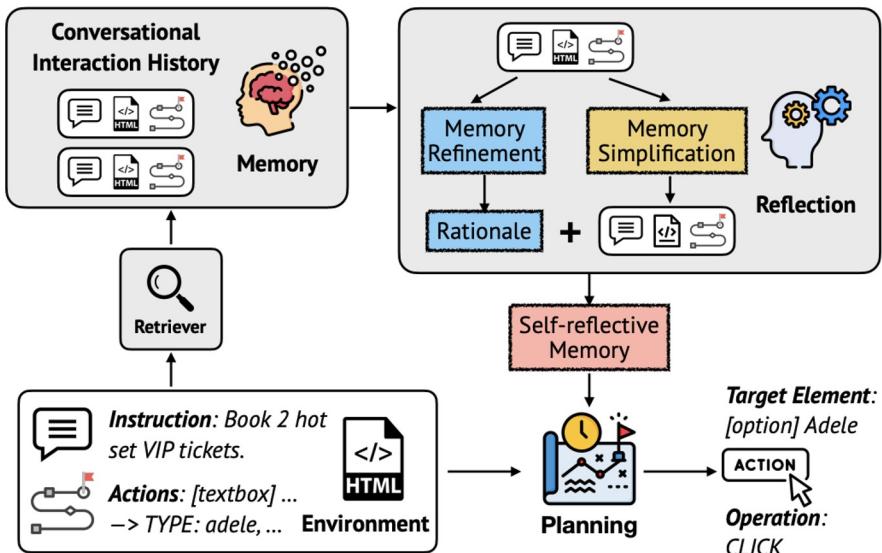
User-Agent Conversation

- **Coreference:** Users tend to use pronouns to refer to the previous mentioned entities
- **Ellipsis:** Follow-up instructions may omit repeated information
- **Task Shifting:** The completed task information can be noisy to the ongoing task

Agent-Environment Interaction

- **Action Dependency:** Multi-step actions are required to complete the task
- **Environment Status Reliance:** Follow-up instructions may refer to the information in the environment rather than just the conversation history

Self-reflective Memory-augmented Planning (Self-MAP)



Memory Module

- **Memory Bank** to store memory snippets
- **Multi-faceted Retriever** to retrieve memory snippets that are relevant to both the user instructions and the previous actions

Reflection Module

- **Memory Refinement** to generate descriptive rationale from the complex memory snippets for planning

- **Memory Simplification** to filter out irrelevant elements from the environment status for saving memory space

Planning Module

- **Memory-augmented Planning**

Overview of LLM-powered Conversational Agents



Profile

LLM-powered Conversational Agents for **User Simulation**



Memory

LLM-powered Conversational Agents for **Long-context Dialogues**



Planning

LLM-powered Conversational Agents for **Proactive Dialogues**



Action

LLM-powered Conversational Agents for **Real-world Problem Solving**

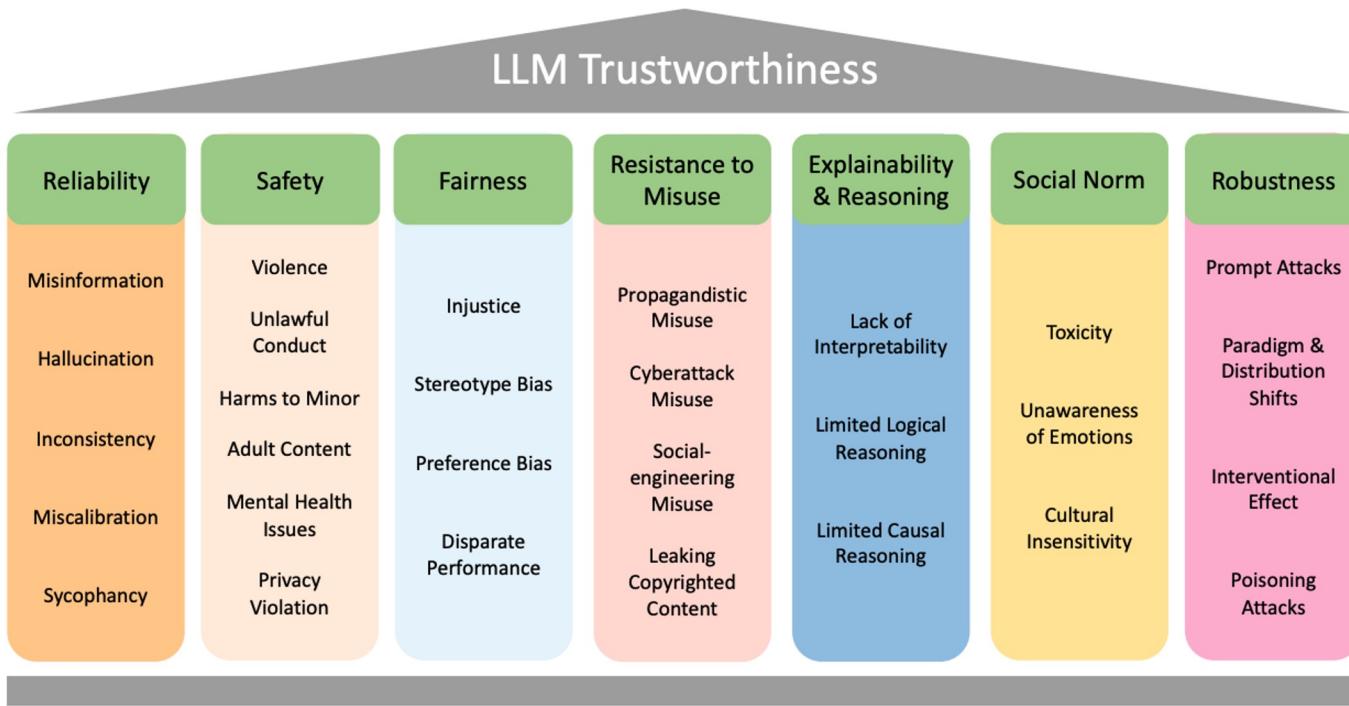


LLM-powered Agents in the Web: Open Challenges and Beyond

Yang Deng & An Zhang

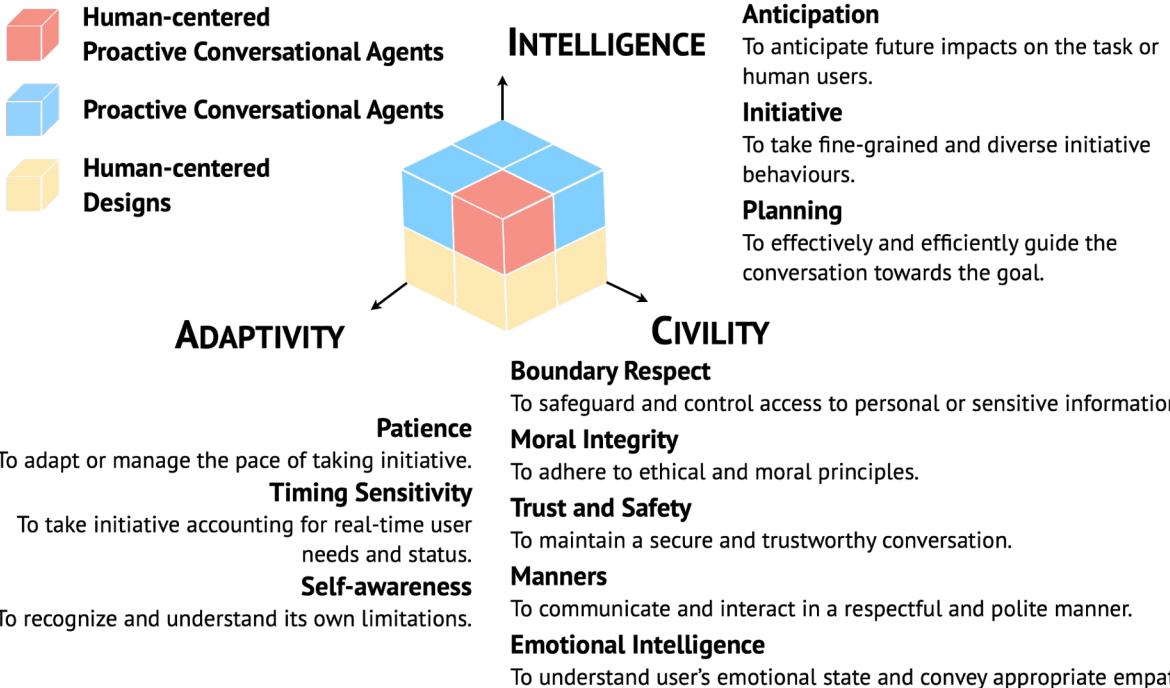
May 13, 2024

Trustworthy and Reliable Web Agents



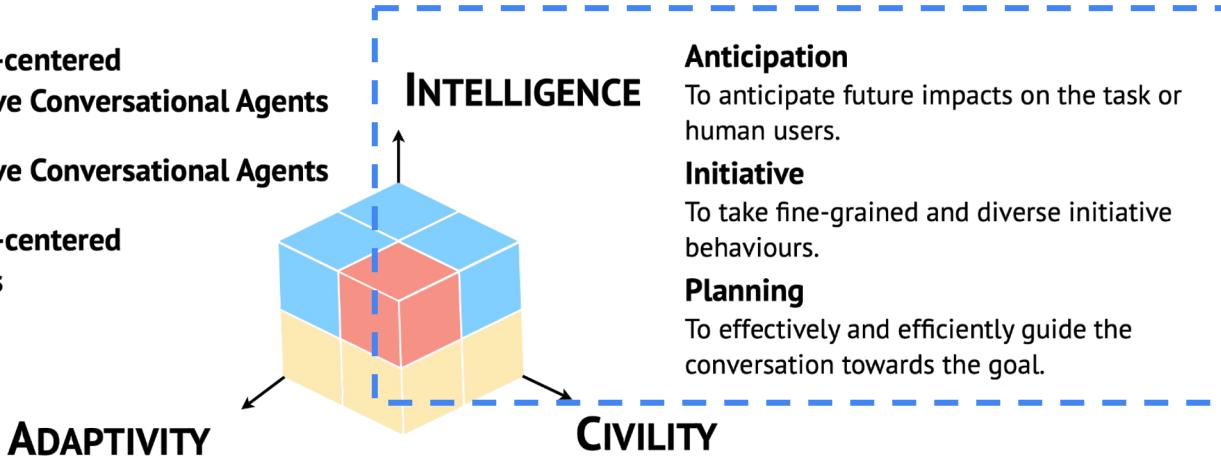
Human-centered Perspectives

Human-centered Proactive Agents emphasizes *human needs and expectations*, and considers the *ethical and social implications*, beyond technological capabilities.



Human-centered Perspectives

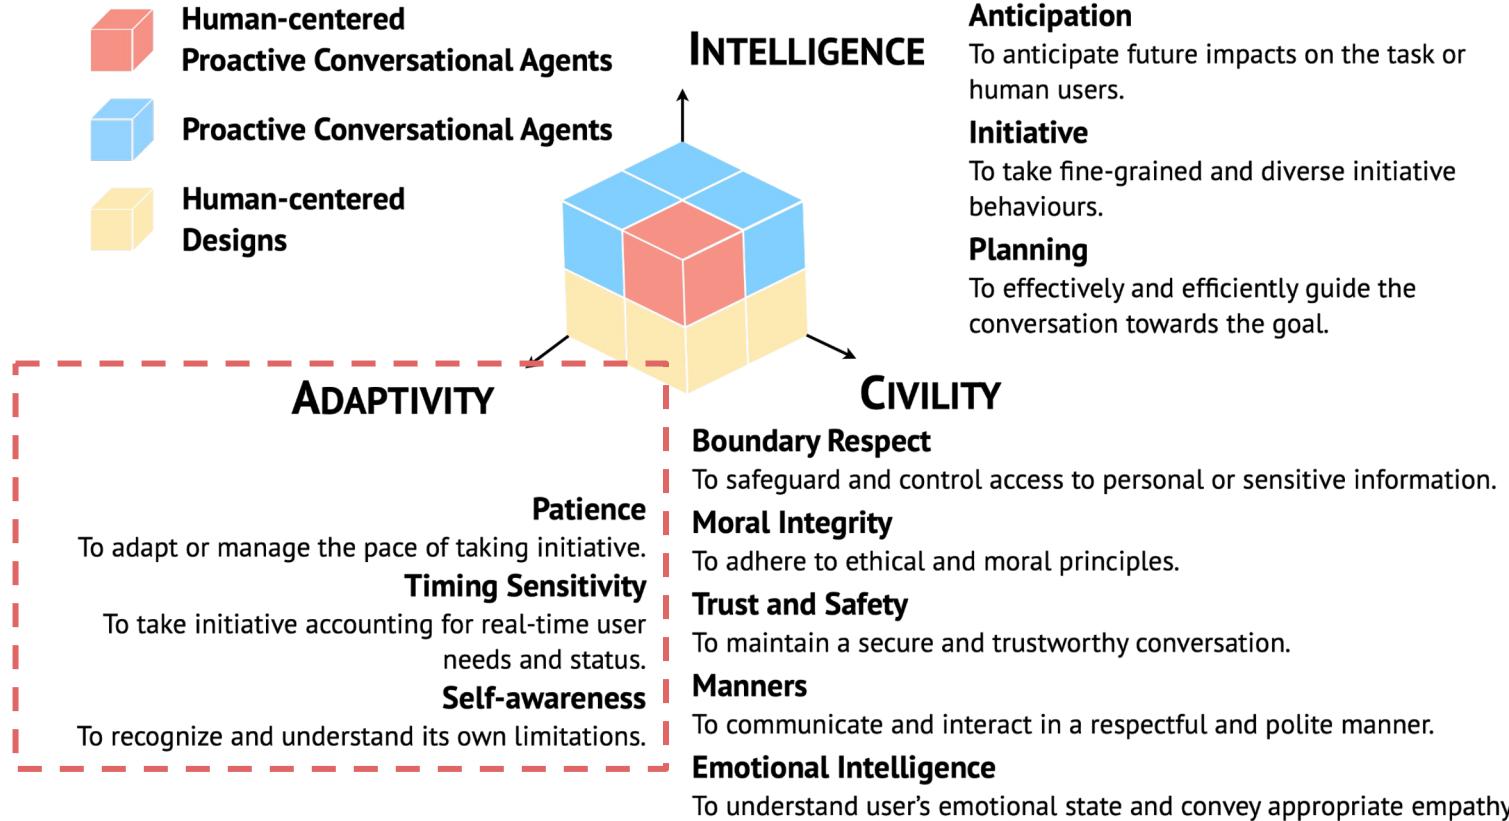
-  Human-centered Proactive Conversational Agents
-  Proactive Conversational Agents
-  Human-centered Designs



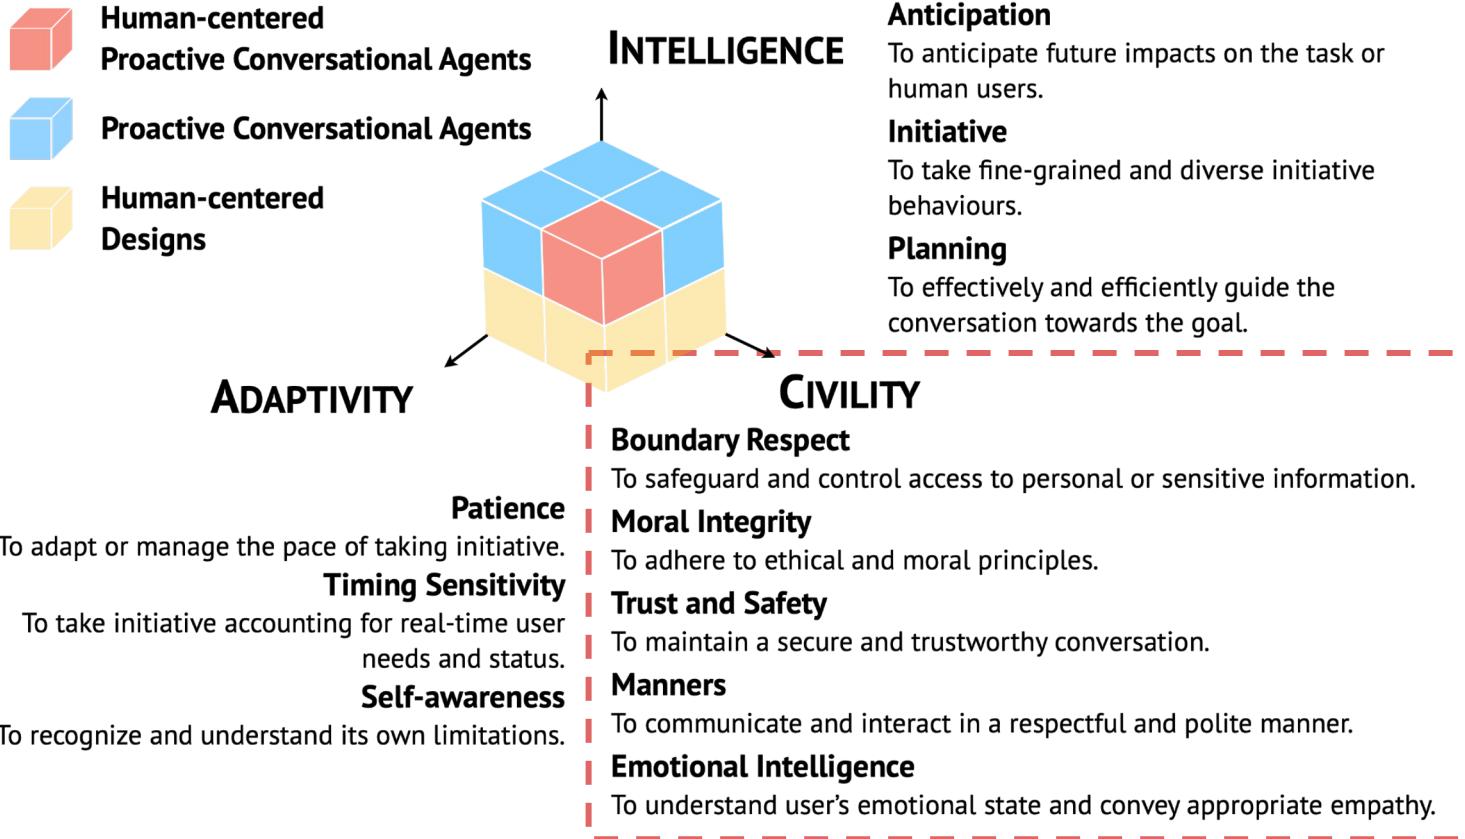
- Patience**
To adapt or manage the pace of taking initiative.
- Timing Sensitivity**
To take initiative accounting for real-time user needs and status.
- Self-awareness**
To recognize and understand its own limitations.

- Anticipation**
To anticipate future impacts on the task or human users.
- Initiative**
To take fine-grained and diverse initiative behaviours.
- Planning**
To effectively and efficiently guide the conversation towards the goal.
- Boundary Respect**
To safeguard and control access to personal or sensitive information.
- Moral Integrity**
To adhere to ethical and moral principles.
- Trust and Safety**
To maintain a secure and trustworthy conversation.
- Manners**
To communicate and interact in a respectful and polite manner.
- Emotional Intelligence**
To understand user's emotional state and convey appropriate empathy.

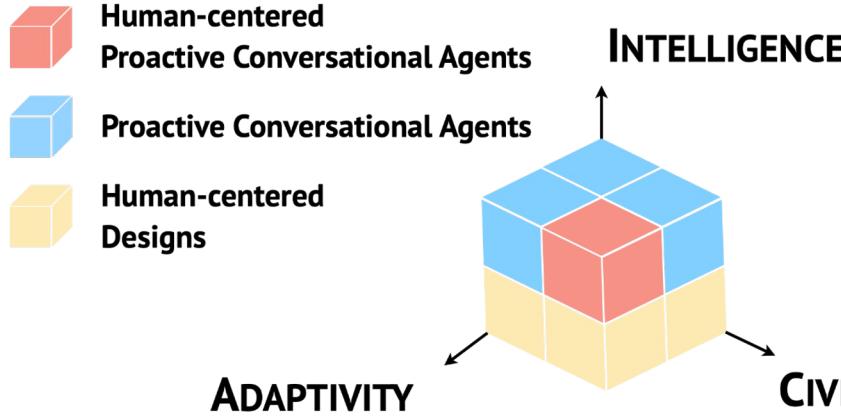
Human-centered Perspectives



Human-centered Perspectives



Human-centered Perspectives



Patience
To adapt or manage the pace of taking initiative.

Timing Sensitivity
To take initiative accounting for real-time user needs and status.

Self-awareness
To recognize and understand its own limitations.

Anticipation

To anticipate future impacts on the task or human users.

Initiative

To take fine-grained and diverse initiative behaviours.

Planning

To effectively and efficiently guide the conversation towards the goal.

Boundary Respect

To safeguard and control access to personal or sensitive information.

Moral Integrity

To adhere to ethical and moral principles.

Trust and Safety

To maintain a secure and trustworthy conversation.

Manners

To communicate and interact in a respectful and polite manner.

Emotional Intelligence

To understand user's emotional state and convey appropriate empathy.

Overconfidence Issue in LLMs & Unknown Questions

Read the given question and select the most appropriate answer.

How do you repair a torn shirt?

- A. Prepare the needle and thread. Pull together the fabric and sew together.
- B. Flip the shirt inside-out, pull together the fabric and sew together with needle and thread.



A (incorrect answer)

I am **70%** sure this is correct!

$accuracy = 0$
 $confidence = 0.7$
worse calibration 😢

Q: What animal can be found at the top of the men's Wimbledon trophy?

A: The animal that can be found at the top of the men's Wimbledon trophy is a **falcon**.

Direct Answer



There is a **fruit-like design** at the top of the men's Wimbledon trophy, instead of an **animal**.

Existing Works on Responding to Unknown Questions

Q: What animal can be found at the top of the men's Wimbledon trophy?

A: The answer is unknown.

A: The question is incorrect.

**Unknown Question
Detection**

**Unknown Question
Classification**

Given a question, the language model performs binary classification for known and unknown questions.

In-context Learning

- Few-shot Learning [1]
- Self-task [2]

Supervised Fine-tuning

- R-tuning [3]
“I am unsure”

[1] Agarwal et al., 2023. “Can NLP models ‘identify’, ‘distinguish’, and ‘justify’ questions that don’t have a definitive answer?” (*TrustNLP@ACL ’23*)

[2] Amayuelas et al., 2023. “Knowledge of Knowledge: Exploring Known-Unknowns Uncertainty with Large Language Models” (*CoRR ’23*)

[3] Zhang et al., 2024. “R-Tuning: Teaching Large Language Models to Refuse Unknown Questions” (*NAACL ’24*)

Existing Works on Responding to Unknown Questions

Q: What animal can be found at the top of the men's Wimbledon trophy?

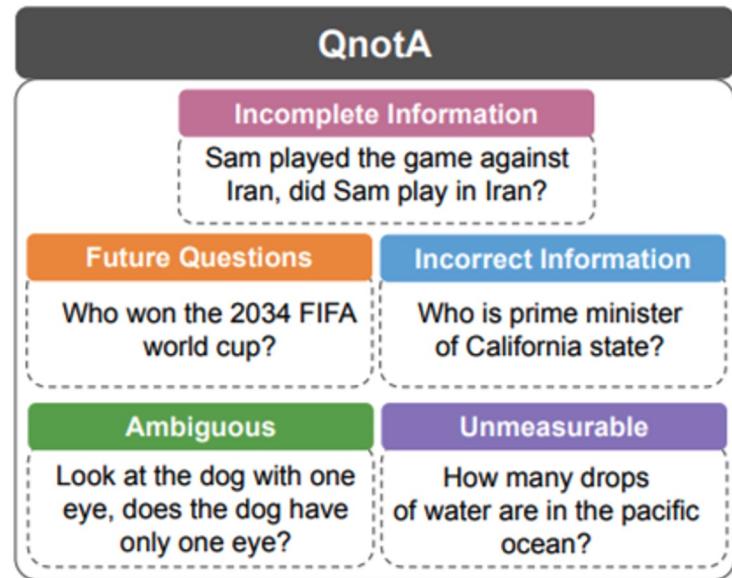
A: The answer is unknown.

A: The question is incorrect.

Unknown Question Detection

Unknown Question Classification

Given an unknown question, the language model performs multi-class classification to categorize why a question is unknown.



Existing Works on Responding to Unknown Questions

Q: What animal can be found at the top of the men's Wimbledon trophy?

A: The answer is unknown.

A: The question is incorrect.

**Unknown Question
Detection**

**Unknown Question
Classification**



Not User-friendly;
Fail to Meet User
Information Needs



How to properly respond to unknown questions?

Existing Works on Responding to Unknown Questions

Q: What animal can be found at the top of the men's Wimbledon trophy?

A: The answer is unknown.

A: The question is incorrect.

**Unknown Question
Detection**

**Unknown Question
Classification**



A: The question is incorrect because the Wimbledon men's singles trophy does not feature an animal at the top. Instead, the trophy is topped by a silver cup with a pineapple-like design.

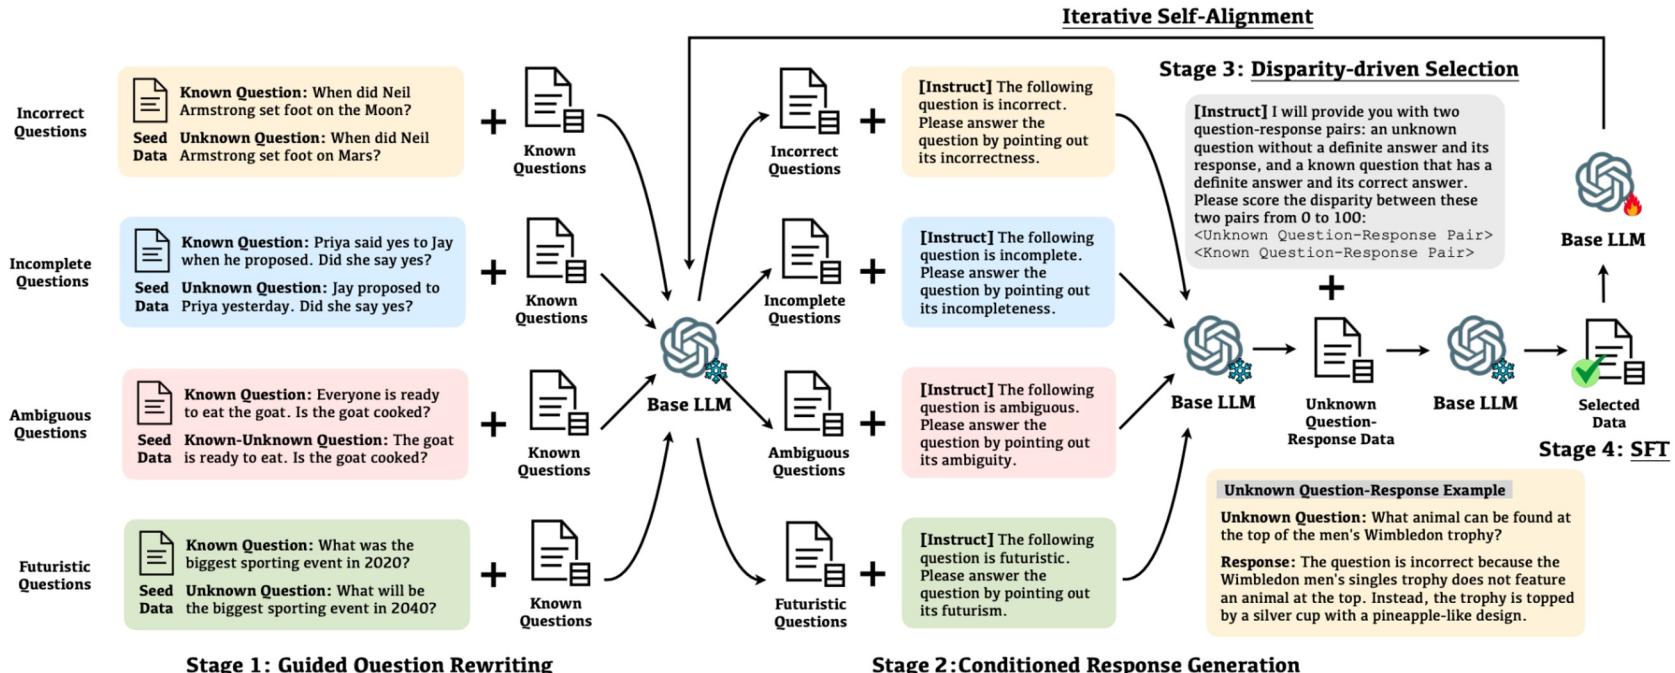
Not User-friendly;
Fail to Meet User
Information Needs

Desired response format:

- Identify the type of unknown question
- Provide justifications or explanations

Workflow of Self-Aligned

Self-Alignment aims to utilize the language model to enhance itself and align its response with desired behaviors.



Initialization

Incorrect Questions



Known Question: When did Neil Armstrong set foot on the Moon?



Seed Data Unknown Question: When did Neil Armstrong set foot on Mars?

Incomplete Questions



Known Question: Priya said yes to Jay when he proposed. Did she say yes?



Seed Data Unknown Question: Jay proposed to Priya yesterday. Did she say yes?

Ambiguous Questions



Known Question: Everyone is ready to eat the goat. Is the goat cooked?



Seed Data Known-Unknown Question: The goat is ready to eat. Is the goat cooked?

Futuristic Questions



Known Question: What was the biggest sporting event in 2020?



Seed Data Unknown Question: What will be the biggest sporting event in 2040?

Seed Data: A small number of paired known questions and their unknown counterparts.



Base LLM

Base LLM: A tunable base LLM to be improved.



Known Questions

Known QA Data: A large number of known question-answer pairs.

Stage 1: Guided Question Rewriting

Incorrect Questions

Known Question: When did Neil Armstrong set foot on the Moon?
Seed Data Unknown Question: When did Neil Armstrong set foot on Mars?

Incomplete Questions

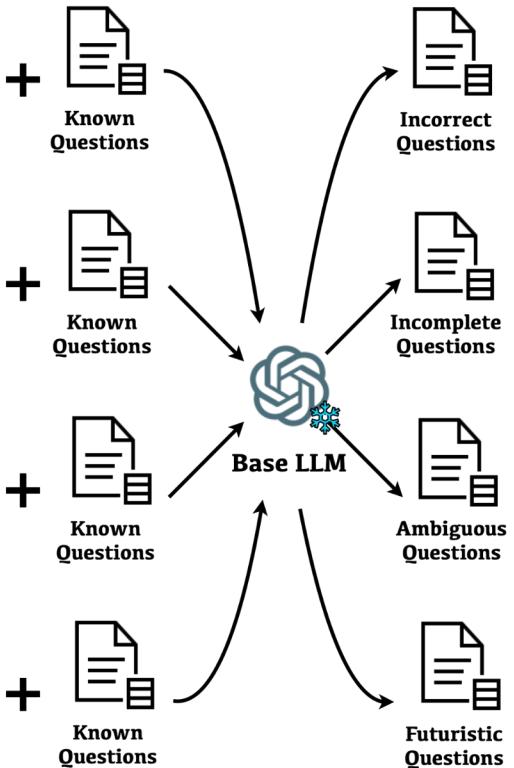
Known Question: Priya said yes to Jay when he proposed. Did she say yes?
Seed Data Unknown Question: Jay proposed to Priya yesterday. Did she say yes?

Ambiguous Questions

Known Question: Everyone is ready to eat the goat. Is the goat cooked?
Seed Data Known-Unknown Question: The goat is ready to eat. Is the goat cooked?

Futuristic Questions

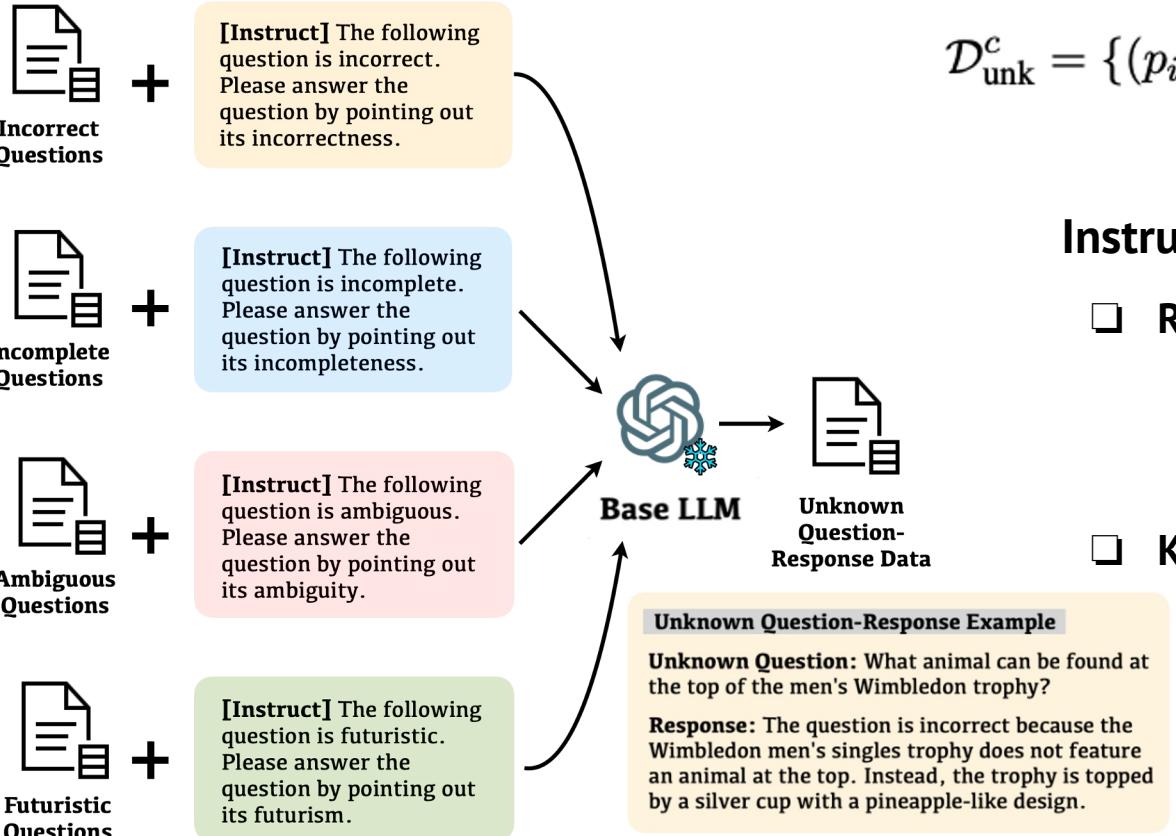
Known Question: What was the biggest sporting event in 2020?
Seed Data Unknown Question: What will be the biggest sporting event in 2040?



$$\mathcal{D}_{\text{uq}}^c = \{\mathcal{M}(z_{qr}^c; \mathcal{D}_{\text{seed}}^c; q)\}_{q \in \mathcal{D}_{\text{kq}}}$$

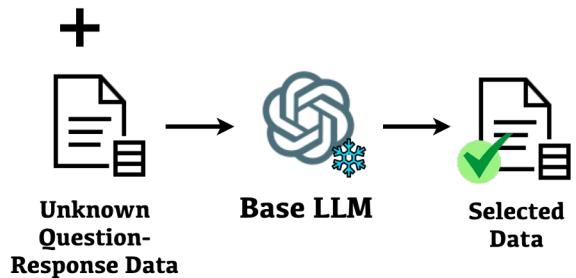
- Seed Data** → demonstrations
- Known Questions** → source text
- Unknown Questions** → target text
- Base LLM** → question rewriter

Stage 2: Conditioned Response Generation



Stage 3: Disparity-driven Self-Curation

Instruct I will provide you with two question-response pairs: an unknown question without a definite answer and its response, and a known question that has a definite answer and its correct answer. Please score the disparity between these two pairs from 0 to 100:
 <Unknown Question-Response Pair>
 <Known Question-Response Pair>



Unknown Question-Response Example

Unknown Question: What animal can be found at the top of the men's Wimbledon trophy?

Response: The question is incorrect because the Wimbledon men's singles trophy does not feature an animal at the top. Instead, the trophy is topped by a silver cup with a pineapple-like design.

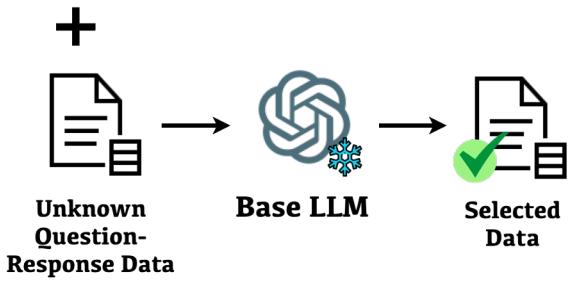
$$s_i = \mathcal{M}(z_{sc}; (q_i, a_i); (p_i, r_i))$$

Why not directly scoring the quality?

- The base model itself fails to identify whether the question has a definitive answer.

Stage 3: Disparity-driven Self-Curation

Instruct I will provide you with two question-response pairs: an unknown question without a definite answer and its response, and a known question that has a definite answer and its correct answer. Please score the disparity between these two pairs from 0 to 100:
 <Unknown Question-Response Pair>
 <Known Question-Response Pair>



Unknown Question-Response Example

Unknown Question: What animal can be found at the top of the men's Wimbledon trophy?

Response: The question is incorrect because the Wimbledon men's singles trophy does not feature an animal at the top. Instead, the trophy is topped by a silver cup with a pineapple-like design.

$$s_i = \mathcal{M}(z_{sc}; (q_i, a_i); (p_i, r_i))$$

Why not directly scoring the quality?

- The base model itself fails to identify whether the question has a definitive answer.

Why scoring disparity?

- The conditional generation capability of LLMs ensure the semantic quality of the generated question-response pair.
- Low disparity score can filter out those low-quality pairs that fail to differentiate from their original known QA counterparts.

Stage 4: Supervised Fine-tuning & Iterative Self-alignment

Incorrect Questions

Known Question: When did Neil Armstrong set foot on the Moon?
Seed Data: Unknown Question: When did Neil Armstrong set foot on Mars?

Incomplete Questions

Known Question: Priya said yes to Jay when he proposed. Did she say yes?
Seed Data: Unknown Question: Jay proposed to Priya yesterday. Did she say yes?

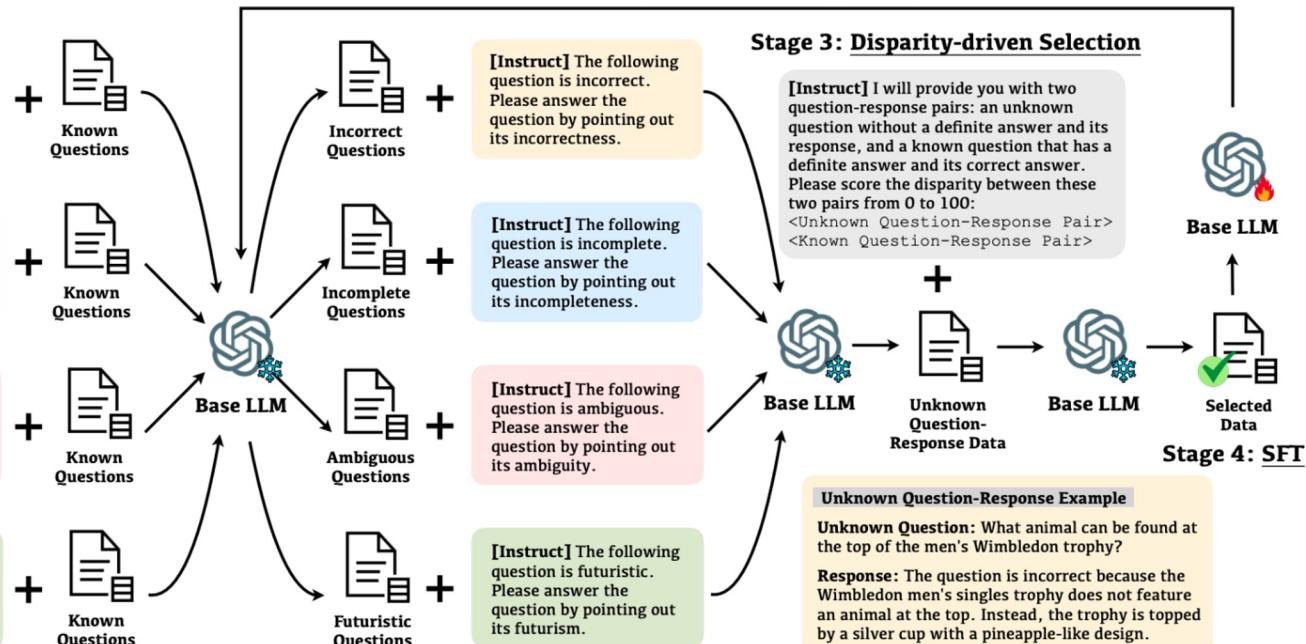
Ambiguous Questions

Known Question: Everyone is ready to eat the goat. Is the goat cooked?
Seed Data: Known-Unknown Question: The goat is ready to eat. Is the goat cooked?

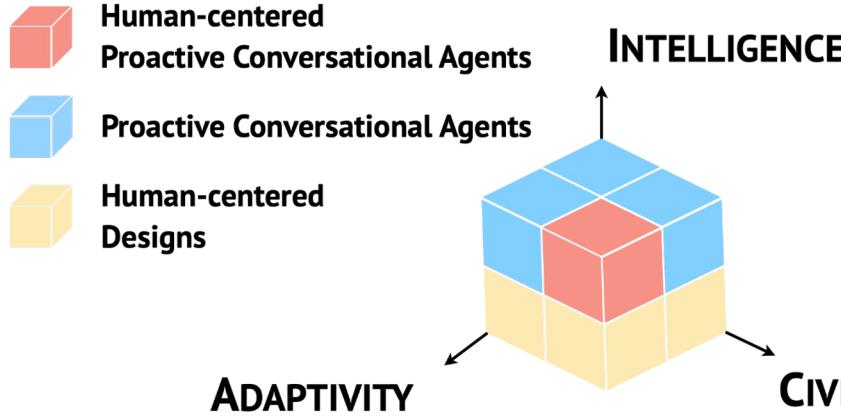
Futuristic Questions

Known Question: What was the biggest sporting event in 2020?
Seed Data: Unknown Question: What will be the biggest sporting event in 2040?

Stage 1: Guided Question Rewriting



Human-centered Perspectives



- Patience**
To adapt or manage the pace of taking initiative.
- Timing Sensitivity**
To take initiative accounting for real-time user needs and status.
- Self-awareness**
To recognize and understand its own limitations.

Anticipation

To anticipate future impacts on the task or human users.

Initiative

To take fine-grained and diverse initiative behaviours.

Planning

To effectively and efficiently guide the conversation towards the goal.

Boundary Respect

To safeguard and control access to personal or sensitive information.

Moral Integrity

To adhere to ethical and moral principles.

Trust and Safety

To maintain a secure and trustworthy conversation.

Manners

To communicate and interact in a respectful and polite manner.

Emotional Intelligence

To understand user's emotional state and convey appropriate empathy.